

Designing Scalable Database ORM

Maghead Database Framework



Yo-An Lin
@c9s

Outline

- Maghead Overview
 - Why? And the history
 - The look of Maghead
 - Performance
- Sharding
 - When do you Shard?
 - What's Sharding?
 - How to Shard?
 - Consistent Hashing
 - Chunks and Chunk Migration
 - Shard Balancing
 - Commands and the APIs

Why **ORM**?

Actually when we
mention **ORM**

Usually it refers to **tools**
including **ORM** and many
other things

**DBAL | ORM | SQL
Builder | Migration |
Database Management**

Some people say ...

DON'T USE **ORM**

Because of
Performance

Use Hand-Written
SQL Queries...

Actually if you

Actually if you

- need to finish 2+ site projects per month
- have different modules shared between different projects with some customization
- need to maintain legacy projects without db query tests
- have frequently schema changes from different team members
- need to remember more than the changes from more than 10, 20, 30 tables

Trust me, You Won't
like it

Writing basic SQL in every
projects is really painful

ORM helps you ...

- Have the domain logics well-organized.
- Perform lightweight automatic migration when you have schema changes.
- Generate the common, basic SQL queries for the most use cases without the need of manually updating the SQL queries.
- Manage the database connections with a centralized simple config file.
- !! Actually you can still do hand-written query if you want

And it saves your **Life**

What's **Maghead**

Maghead Database Framework

Object Relation Mapper

DataBase Abstraction Layer

SQL Builder

Connection Manager

Database Manager

Table Definition Parser

Automatic Migration

Migration Manager

And many components

The components are
designed to be used
separately.

The project was started
since 2010

7 years

The project was created to
have the **dynamic schema**
in different projects

How does it looks like?

YAML configuration

cli:

bootstrap: vendor/autoload.php

schema:

auto_id: true

finders:

- { name: ComposerSchemaFinder, args: ["composer.json"] }

...

instance:

 local:

 dsn: 'mysql:host=localhost'

 user: root

 driver: mysql

 host: localhost

 password: null

 query_options: { }

 connection_options:

 1002: 'SET NAMES utf8'

...

...

databases:

 master:

 dsn: 'mysql:host=localhost;dbname=testing'

 host: localhost

 user: root

 driver: mysql

 database: testing

 password: null

 query_options: { }

 connection_options:

 1002: 'SET NAMES utf8'

Simple bootstrapping

Bootstrap Code

```
require 'vendor/autoload.php';

use Maghead\Runtime\Config\FileConfigLoader;
use Maghead\Runtime\Bootstrap;

$config = FileConfigLoader::load( __DIR__ . '/db/config/
database.yml');
Bootstrap::setup($config);
```

Bootstrap Code

```
require 'vendor/autoload.php';  
  
use Maghead\Runtime\Config\FileConfigLoader;  
use Maghead\Runtime\Bootstrap;  
  
$config = FileConfigLoader::load('db/config/database.yml');  
Bootstrap::setup($config);
```



2 lines only!

ActiveRecord pattern

Most of **ActiveRecord**
patterns are implemented
in this way

```
$record = new User;  
$record->account = "c9s";  
$record->password = sha1("mypassword");  
$record->save();
```

Different States:

1. existing record
2. non-existing record (before inserting the row)

```
function updatePassword($user) {  
    // check the existence because the object  
    // might be not able to update the  
    // password  
    if (!$user->hasKey()) {  
        throw new LogicException;  
    }  
    // update ...  
}
```

```
function updatePassword($user) {
```

```
    // check the existence because the object
```

```
    // might be not able
```

The user record might be not created yet.

```
    // password
```

```
    if (!$user->hasKey()) {
```

```
        throw new LogicException;
```

```
    }
```

```
    // update ...
```

```
}
```

```
function updatePassword($user) {  
    // check the existence because the object  
    // might be not able to update the  
    // password  
    if (!$user->hasKey()) {  
        throw new LogicException:  
    }  
    // update ...  
}
```

Check the key to verify the existence

The **Maghead** Way

```
$record = User::load([ "account" => "timcook" ]);
```

If the returned \$record is not false, then it must be an existing row.

```
$user = User::createAndLoad([ "account" => "timcook" ]);
```

If the returned \$record is not false, then it must be an existing row.


```
$user->update([ "password" => sha1("newpw") ]);
```

Safely Update the row without the concern

Repository pattern


```
$repo = new ProductRepo($write, $read);
```

```
$repo = Product::repo('master', 'slave');
```

```
$repo = Product::masterRepo();
```

Dynamic **Schema**

```
namespace Todos\Model;  
  
use Maghead\Schema\DeclareSchema;  
  
class TodoSchema extends DeclareSchema  
{  
    public function schema()  
    {  
  
    }  
}  
}
```



Column Definition Goes Here

```
namespace Todos\Model;

use Maghead\Schema\DeclareSchema;

class TodoSchema extends DeclareSchema
{
    public function schema()
    {
        $this->column('title')
            ->varchar(128)
            ->required()
            ;
    }
}
```



```
$this->column('done')  
->boolean()  
->default(false);
```

```
$this->column('description')  
->text();
```

```
$this->column('created_at')  
->timestamp()  
->default(function() {  
    return date('c');  
});
```



Dynamic Default Value

```
$this->column('extra')  
->text()  
->isa('json');
```



Automatically Define JSON
Inflator and Deflator

```
$this->column('address')  
->varchar(64)  
->validator(function ($val, $args) {  
    if (strlen($val) < 6) {  
        return [false, "Invalid address"];  
    }  
    return [true, "Good address"];  
})
```



Validation Integration

```
$this->column('id', 'ai-pk');  
    // AutoIncrementPrimaryKeyColumn
```


```
$this->column('id', 'uuid-pk');  
    // UUIDPrimaryKeyColumn
```

And can be compiled into
the static schema class

- Maghead\Schema\DeclareSchema

declarative implementation

- Maghead\Schema\DeclareColumn



Compile

- Maghead\Schema\RuntimeSchema

lightweight implementation

- Maghead\Schema\RuntimeColumn

maghead schema build

```
namespace Todos\Model;

class TodoSchemaProxy
    extends RuntimeSchema
{

    const SCHEMA_CLASS = 'Todos\\Model\\TodoSchema';

    const LABEL = 'Todo';

    const MODEL_NAME = 'Todo';

    const MODEL_NAMESPACE = 'Todos\\Model';

    const MODEL_CLASS = 'Todos\\Model\\Todo';

    const REPO_CLASS = 'Todos\\Model\\TodoRepoBase';

    const COLLECTION_CLASS = 'Todos\\Model\\TodoCollection';
```

... •

And This Reduces The
Runtime Overhead

Pre-Generated SQL Queries

```
class ProductRepoBase extends Repo {

  const FIND_BY_PRIMARY_KEY_SQL = 'SELECT * FROM products WHERE id = ? LIMIT 1';

  const DELETE_BY_PRIMARY_KEY_SQL = 'DELETE FROM products WHERE id = ?';

  const FETCH_CREATED_BY_SQL = 'SELECT * FROM users WHERE id = ? LIMIT 1';

  const FETCH_UPDATED_BY_SQL = 'SELECT * FROM users WHERE id = ? LIMIT 1';

  const FETCH_PRODUCT_FEATURES_SQL = 'SELECT * FROM product_feature_junction WHERE product_id = ?';

  const FETCH_PRODUCT_PRODUCTS_SQL = 'SELECT * FROM product_products WHERE product_id = ?';

  const FETCH_IMAGES_SQL = 'SELECT * FROM product_images WHERE product_id = ?';

  ... •
```

Pre-Generated Query Methods

```
class ProductRepoBase
    extends Repo
{
    public function fetchCategoryOf(Model $record)
    {
        if (!$this->fetchCategoryStm) {
            $this->fetchCategoryStm =
                $this->read->prepare(self::FETCH_CATEGORY_SQL);
            $this->fetchCategoryStm->setFetchMode(
                PDO::FETCH_CLASS,
                \ProductBundle\Model\Category::class, [$this]);
        }
        $this->fetchCategoryStm->execute([$record->category_id]);
        $obj = $this->fetchCategoryStm->fetch();
        $this->fetchCategoryStm->closeCursor();
        return $obj;
    }
}
```



```
class ProductBase
  extends Model
{
  public function getProductTags()
  {
    $collection = new \ProductBundle\Model\ProductTagCollection;
    $collection->where()->equal("product_id", $this->id);
    $collection->setPresetVars([ "product_id" => $this->id ]);
    return $collection;
  }
}
```

```
class ProductBase
    extends Model
{
    public function getTags()
    {
        $collection = new \ProductBundle\Model\TagCollection;
        $collection->joinTable('product_tag_junction', 'j', 'INNER')
            ->on("j.tag_id = {$collection->getAlias()}.id");
        $collection->where()->equal('j.product_id', $this->id);
        $parent = $this;
        $collection->setAfterCreate(function($record, $args) use ($parent) {
            $a = [
                'tag_id' => $record->get("id"),
                'product_id' => $parent->id,
            ];
            if (isset($args['product_tags'])) {
                $a = array_merge($args['product_tags'], $a);
            }
            return \ProductBundle\Model\ProductTag::createAndLoad($a);
        });
        return $collection;
    }
}
```

Statements are prepared
and cached in the
Repository of each **Model**

Same as

your hand-written queries

Building **Table Schema**

```
maghead sql --rebuild node1
```

**One Schema to Rule
Them All**

```
$create = $product->asCreateAction();
```



```
$create = $product->asCreateAction();
```

```
{{ RecordAction.renderSignatureWidget | raw }}  
{{ RecordAction.renderCSRFTokenWidget | raw }}
```

```
{% if Record.hasKey %}  
    {{RecordAction.renderKeyWidget | raw}}  
{% endif %}
```

```
{{RecordAction.renderField('account') | raw }}  
{{RecordAction.renderField('password') | raw }}
```



Render Web Form Fields

Relationships

```
$this->belongsTo('book', BookSchema::class, 'id', 'book_id')  
->onDelete('CASCADE')  
->onUpdate('CASCADE')  
;
```

```
$this->many('author_books',  
    AuthorBookSchema::class, 'author_id', 'id');
```

```
$this->manyToMany('books', 'author_books', 'book');
```

```
$this->many('author_books',  
    AuthorBookSchema::class, 'author_id', 'id');  
  
$this->manyToMany('books', 'author_books', 'book');  
  
    foreach ($author->books as $book) {  
        ...  
    }
```

Database Management Commands

maghead db create master

maghead db recreate master

maghead db create node1

maghead db drop node1

Repository pattern for **Multiple Database** Connections

The **Doctrine** Way

```
use Doctrine\DBAL\DriverManager;
$conn = DriverManager::getConnection(array(
    'wrapperClass' =>
'Doctrine\DBAL\Connections\MasterSlaveConnection',
    'driver' => 'pdo_mysql',
    'keepSlave' => true,
    'master' => array(
        'user' => 'ideato',
        'password' => 'ideato',
        'dbname' => 'db_ideato'
    ),
    'slaves' => array(
        array(
            'user' => 'ideato',
            'password' => 'ideato',
            'dbname' => 'db_ideato_slave'
        )
    )
);
$entityManager = EntityManager::create($conn, $config);
```

```
// https://gist.github.com/ricfrank/d6f6317a1a1434cdc364  
$entityManager = EntityManager::create($conn, $config);  
$productRepository = $entityManager->getRepository('Product');  
  
$masterSlaveConn->connect('slave');  
$productRepository = $entityManager->getRepository('Product');  
  
$masterSlaveConn->connect('master');  
$productRepository = $entityManager->getRepository('Product');
```

```
$entityManager = EntityManager::create($conn, $config);  
$productRepository = $entityManager->getRepository('Product');  
  
$masterSlaveConn->connect('slave');
```

```
$product = new Product();  
$product->setName("nuovo_prod1");  
$product->setCategory("oeoeoeoe");  
$entityManager->persist($product);  
$entityManager->flush();
```

```
$productRepository = $entityManager->getRepository('Product');  
  
$masterSlaveConn->connect('master');  
$productRepository = $entityManager->getRepository('Product');
```

The **Maghead** Way


```
$ret = Product::repo("master")  
    ->create([ "name" => "Samsung S3" ]);
```

```
$ret = Product::masterRepo()  
    ->create([ "name" => "Samsung S3" ]);
```

```
$ret = Product::repo("node2")  
      ->create([ "name" => "iPhone 7" ]);
```

```
$ret = Product::repo(new Connection("mysql:host=db1"))  
    ->create([ "name" => "Samsung S3" ]);
```

Automatic Migration

Known as lightweight migration in iOS app development

maghead diff

Performing comparison...

+ table 'metric_values' examples/metric/Model/MetricValueSchema.php

A val	double
A published_at	timestamp
A unit	varchar

+ table 'Edm' tests/apps/TestApp/Model/EdmSchema.php

D id	int
A edmNo	int
A edmTitle	varchar
A edmStart	date
A edmEnd	date
A edmContent	text
A edmUpdatedOn	timestamp

+ table 'i_d_numbers' tests/apps/TestApp/Model/IDNumberSchema.php

A id_number	varchar
-------------	---------

+ table 'posts' tests/apps/TestApp/Model/PostSchema.php

A title	varchar
A content	text
A status	varchar
A created_at	timestamp
A created_by	int

+ table 'tables' tests/apps/TestApp/Model/TableSchema.php

M id	int
A title	varchar
A columns	text
A rows	text

Performing comparison...

+ table 'metric_values' examples/metric/Model/MetricValueSchema.php

A val double

A published_at timestamp

A unit varchar

+ table 'Edm' tests/apps/TestApp/Model/EdmSchema.php

D id int

D: 欄位刪除

A edmNo int

A edmTitle varchar

A edmStart date

A edmEnd date

A edmContent text

A edmUpdatedOn timestamp

+ table 'i_d_numbers' tests/apps/TestApp/Model/IDNumberSchema.php

A id_number varchar

+ table 'posts' tests/apps/TestApp/Model/PostSchema.php

A title varchar

A content text

A status varchar

A created_at timestamp

A created_by int

+ table 'tables' tests/apps/TestApp/Model/TableSchema.php

M id int

A title varchar

A columns text

A rows text

Performing comparison...

+ table 'metric_values' examples/metric/Model/MetricValueSchema.php

A val double

A published_at timestamp

A unit varchar

+ table 'Edm' tests/apps/TestApp/Model/EdmSchema.php

D id int

A edmNo int

A edmTitle varchar

A edmStart date

A edmEnd date

A edmContent text

A edmUpdatedOn timestamp

+ table 'i_d_numbers' tests/apps/TestApp/Model/IDNumberSchema.php

A id_number varchar

+ table 'posts' tests/apps/TestApp/Model/PostSchema.php

A title varchar

A content text

A status varchar

A created_at timestamp

A created_by int

+ table 'tables' tests/apps/TestApp/Model/TableSchema.php

M id int

A title varchar

A columns text

A rows text

A: 欄位新增

Performing comparison...

+ table 'metric_values' examples/metric/Model/MetricValueSchema.php

A val double
A published_at timestamp
A unit varchar

+ table 'Edm' tests/apps/TestApp/Model/EdmSchema.php

D id int
A edmNo int
A edmTitle varchar
A edmStart date
A edmEnd date
A edmContent text
A edmUpdatedOn timestamp

+ table 'i_d_numbers' tests/apps/TestApp/Model/IDNumberSchema.php

A id_number varchar

+ table 'posts' tests/apps/TestApp/Model/PostSchema.php

A title varchar
A content text
A status varchar
A created_at timestamp
A created_by int

+ table 'tables' tests/apps/TestApp/Model/TableSchema.php

M id int
A title varchar
A columns text
A rows text

M: 欄位修改

maghead m auto

Performing automatic upgrade over data source: master

Begining transaction...

Performing Query: ALTER TABLE `addresses` DROP FOREIGN KEY `addresses_ibfk_1`

Performing Query: ALTER TABLE `books` DROP FOREIGN KEY `books_ibfk_1`

Performing Query: ALTER TABLE `books` DROP FOREIGN KEY `books_ibfk_2`

Performing Query: ALTER TABLE `author_books` DROP FOREIGN KEY `author_books_ibfk_1`

Performing Query: ALTER TABLE `author_books` DROP FOREIGN KEY `author_books_ibfk_2`

Performing Query: ALTER TABLE `book_tags` DROP FOREIGN KEY `book_tags_ibfk_1`

Importing schema: MetricApp\Model\MetricValueSchema

Performing Query:


```
CREATE TABLE IF NOT EXISTS `metric_values` (  
  `id` int UNSIGNED NOT NULL PRIMARY KEY AUTO_INCREMENT,  
  `val` double(5,3) NOT NULL DEFAULT 0,  
  `published_at` timestamp NOT NULL,  
  `unit` varchar(3)  
) ENGINE=InnoDB;
```

全自動產生 Alter Table Query

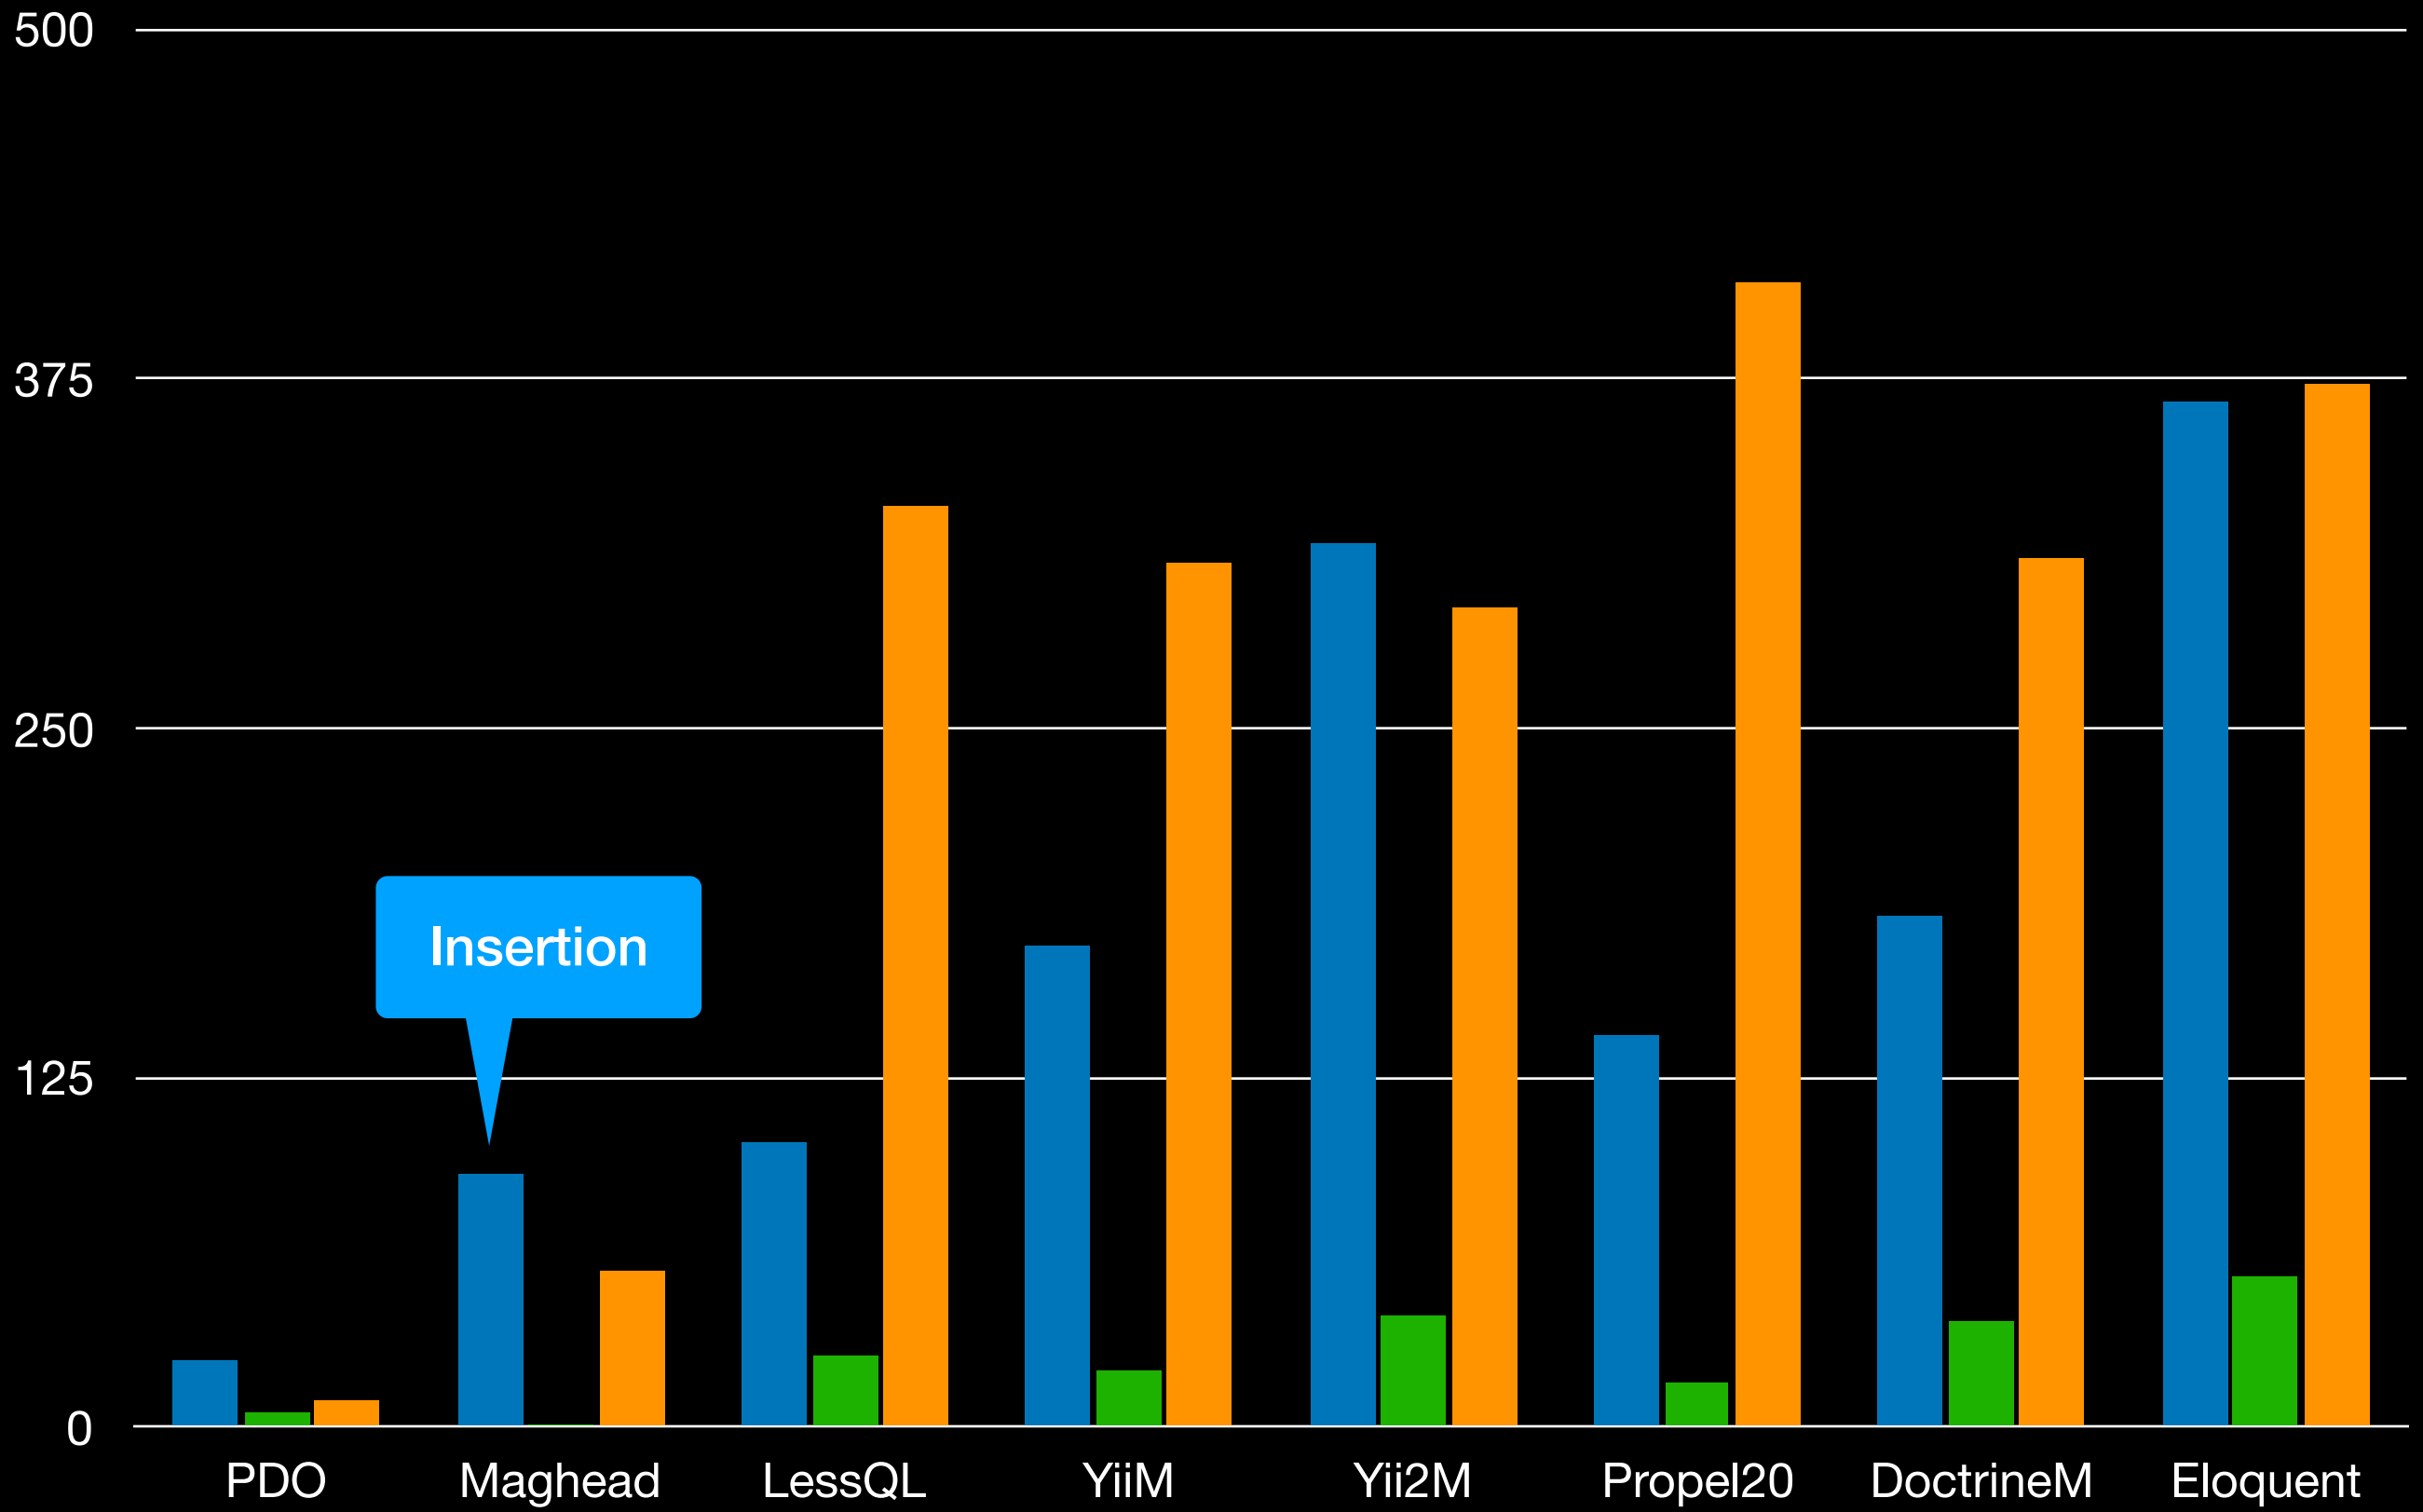
Database Support


SQLite
MySQL
PostgreSQL

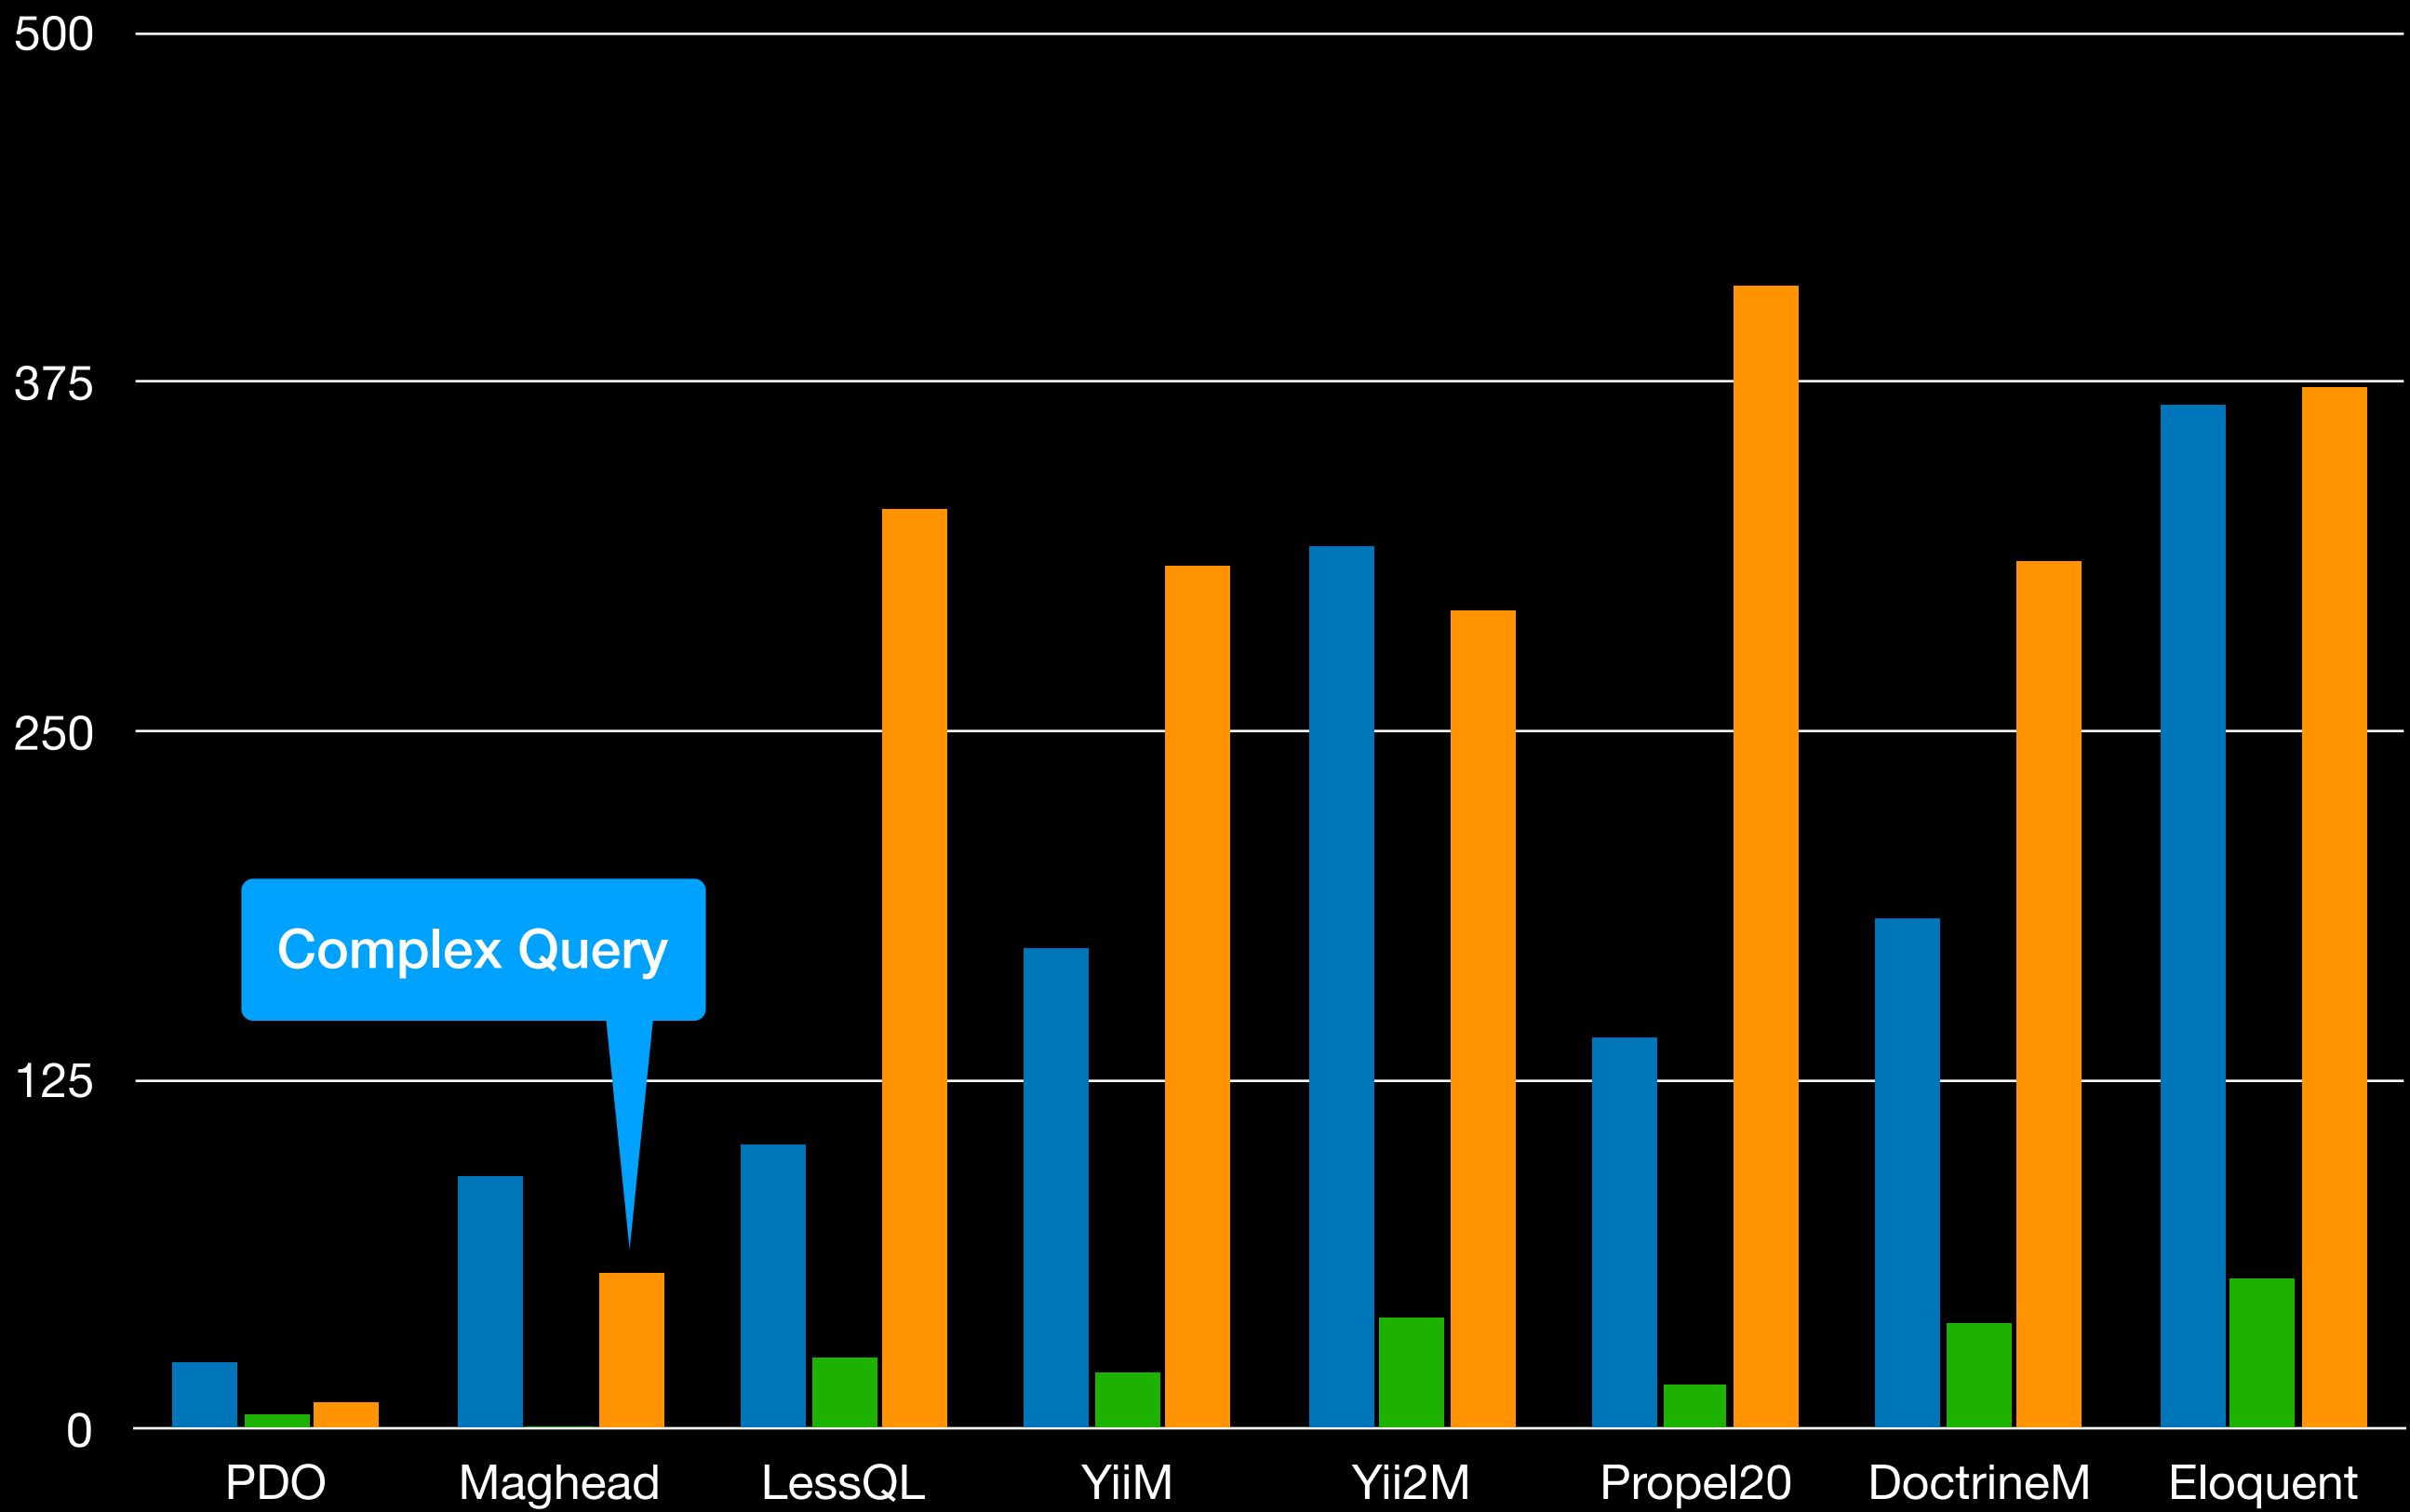
How Fast?




Benchmark: score -> lower is better



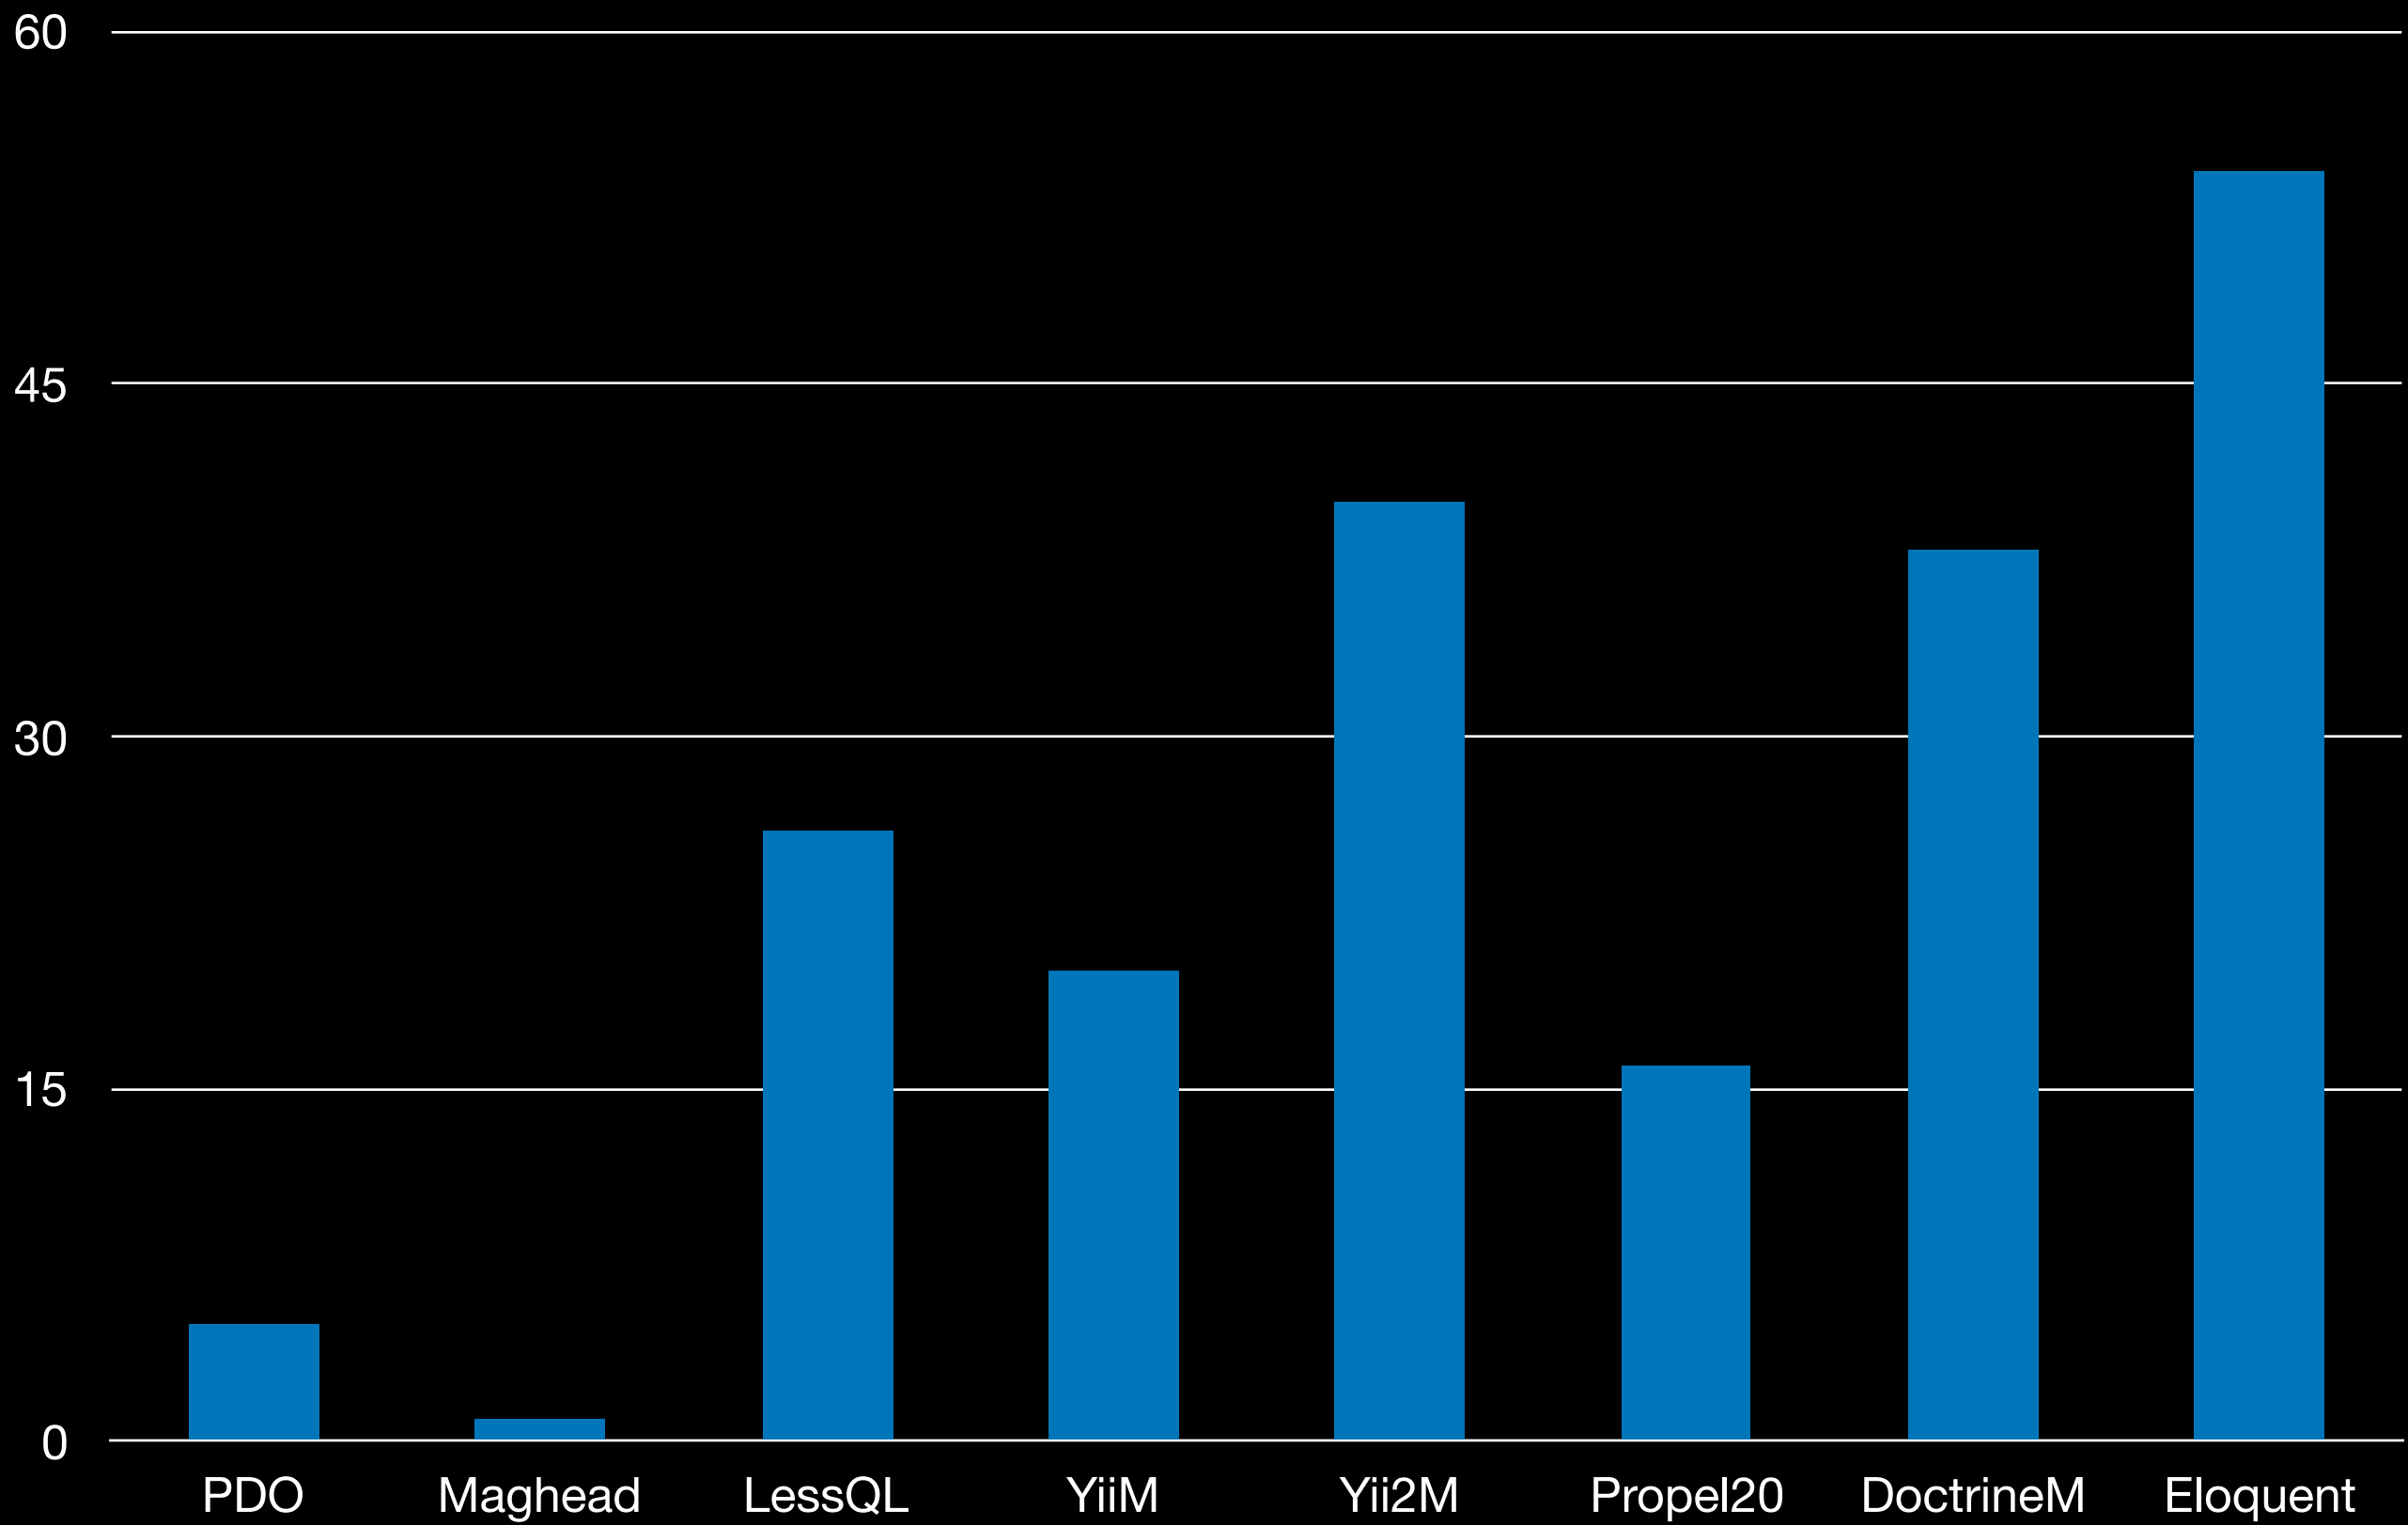






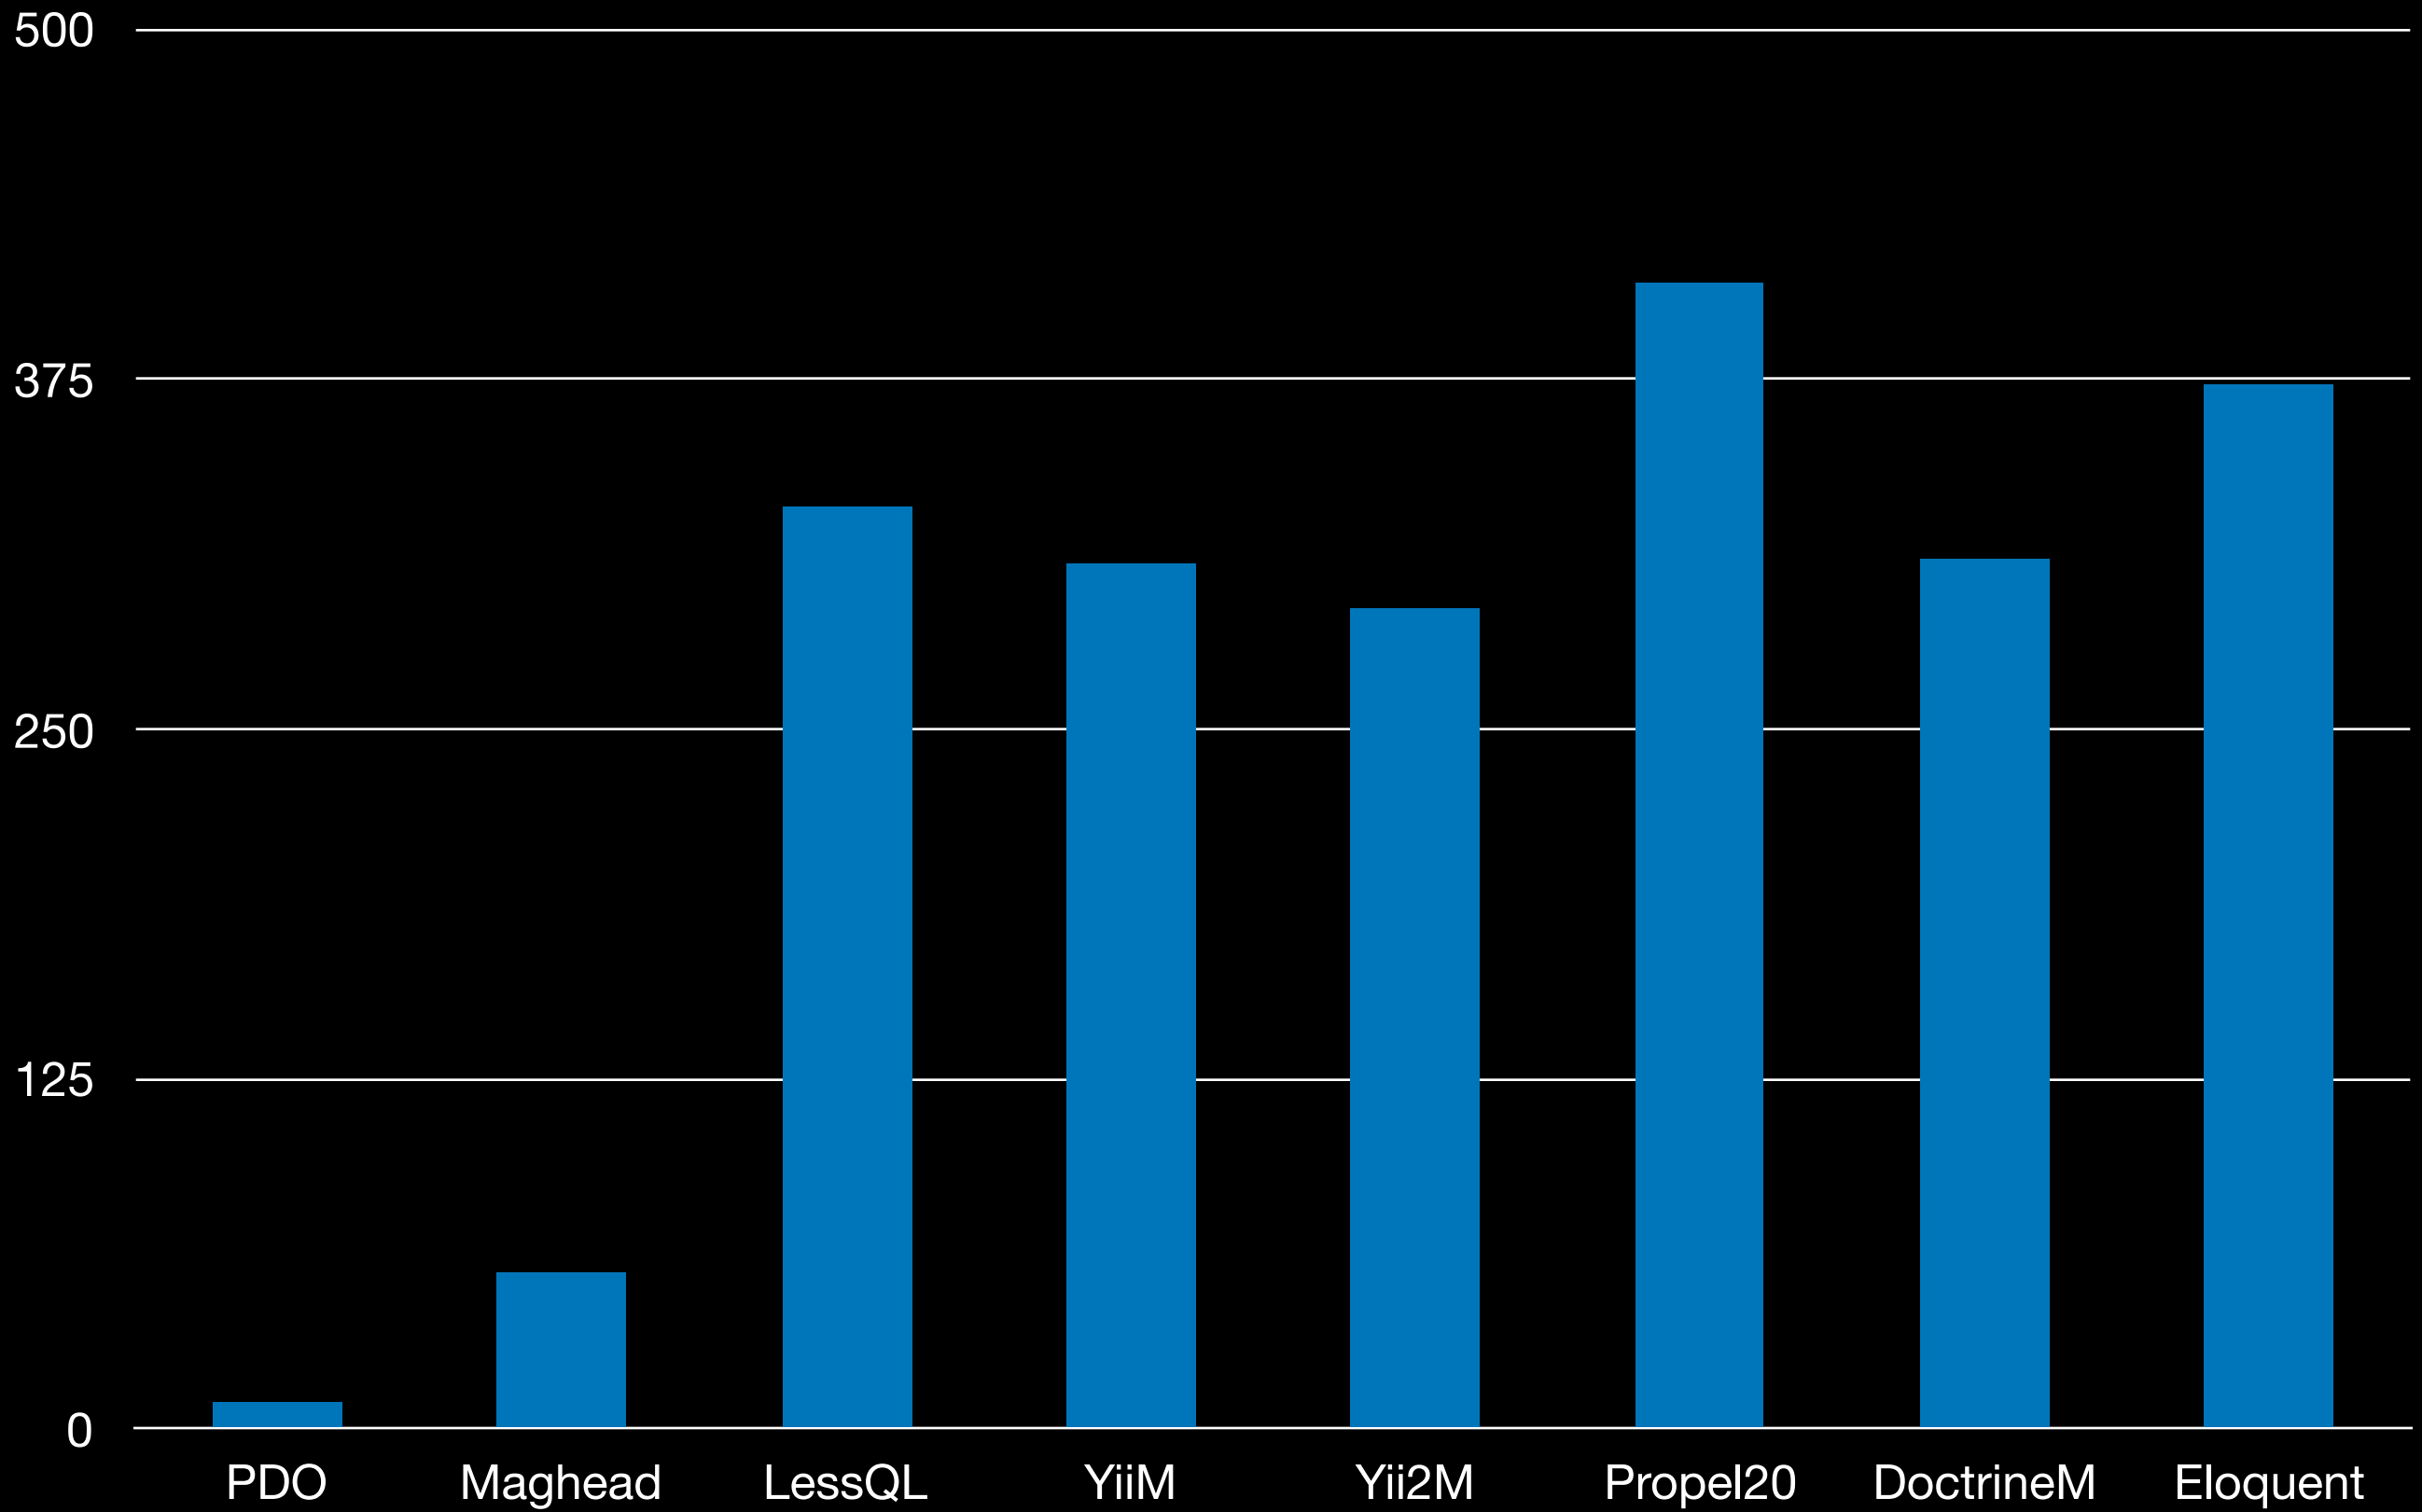
Benchmark: Insert

lower is better




Benchmark: Find PK

lower is better



Benchmark: Complex Query

lower is better



The fastest pure PHP
ORM implementation

Sharding

When do you
Shard?


You started from one
single database server

When rows grows up
to **million...**

queries become
slow ...

more read queries start
hitting on the database

then you split **read/write**
to
master/replicas



finally the write load
increases, **master** is
crying

And rows **grows up** to
10M+...

Covering index doesn't
work well anymore

Any query without
covering index won't
come back in 2-5 minutes

What **NOW**?

Upgrade the hardware?

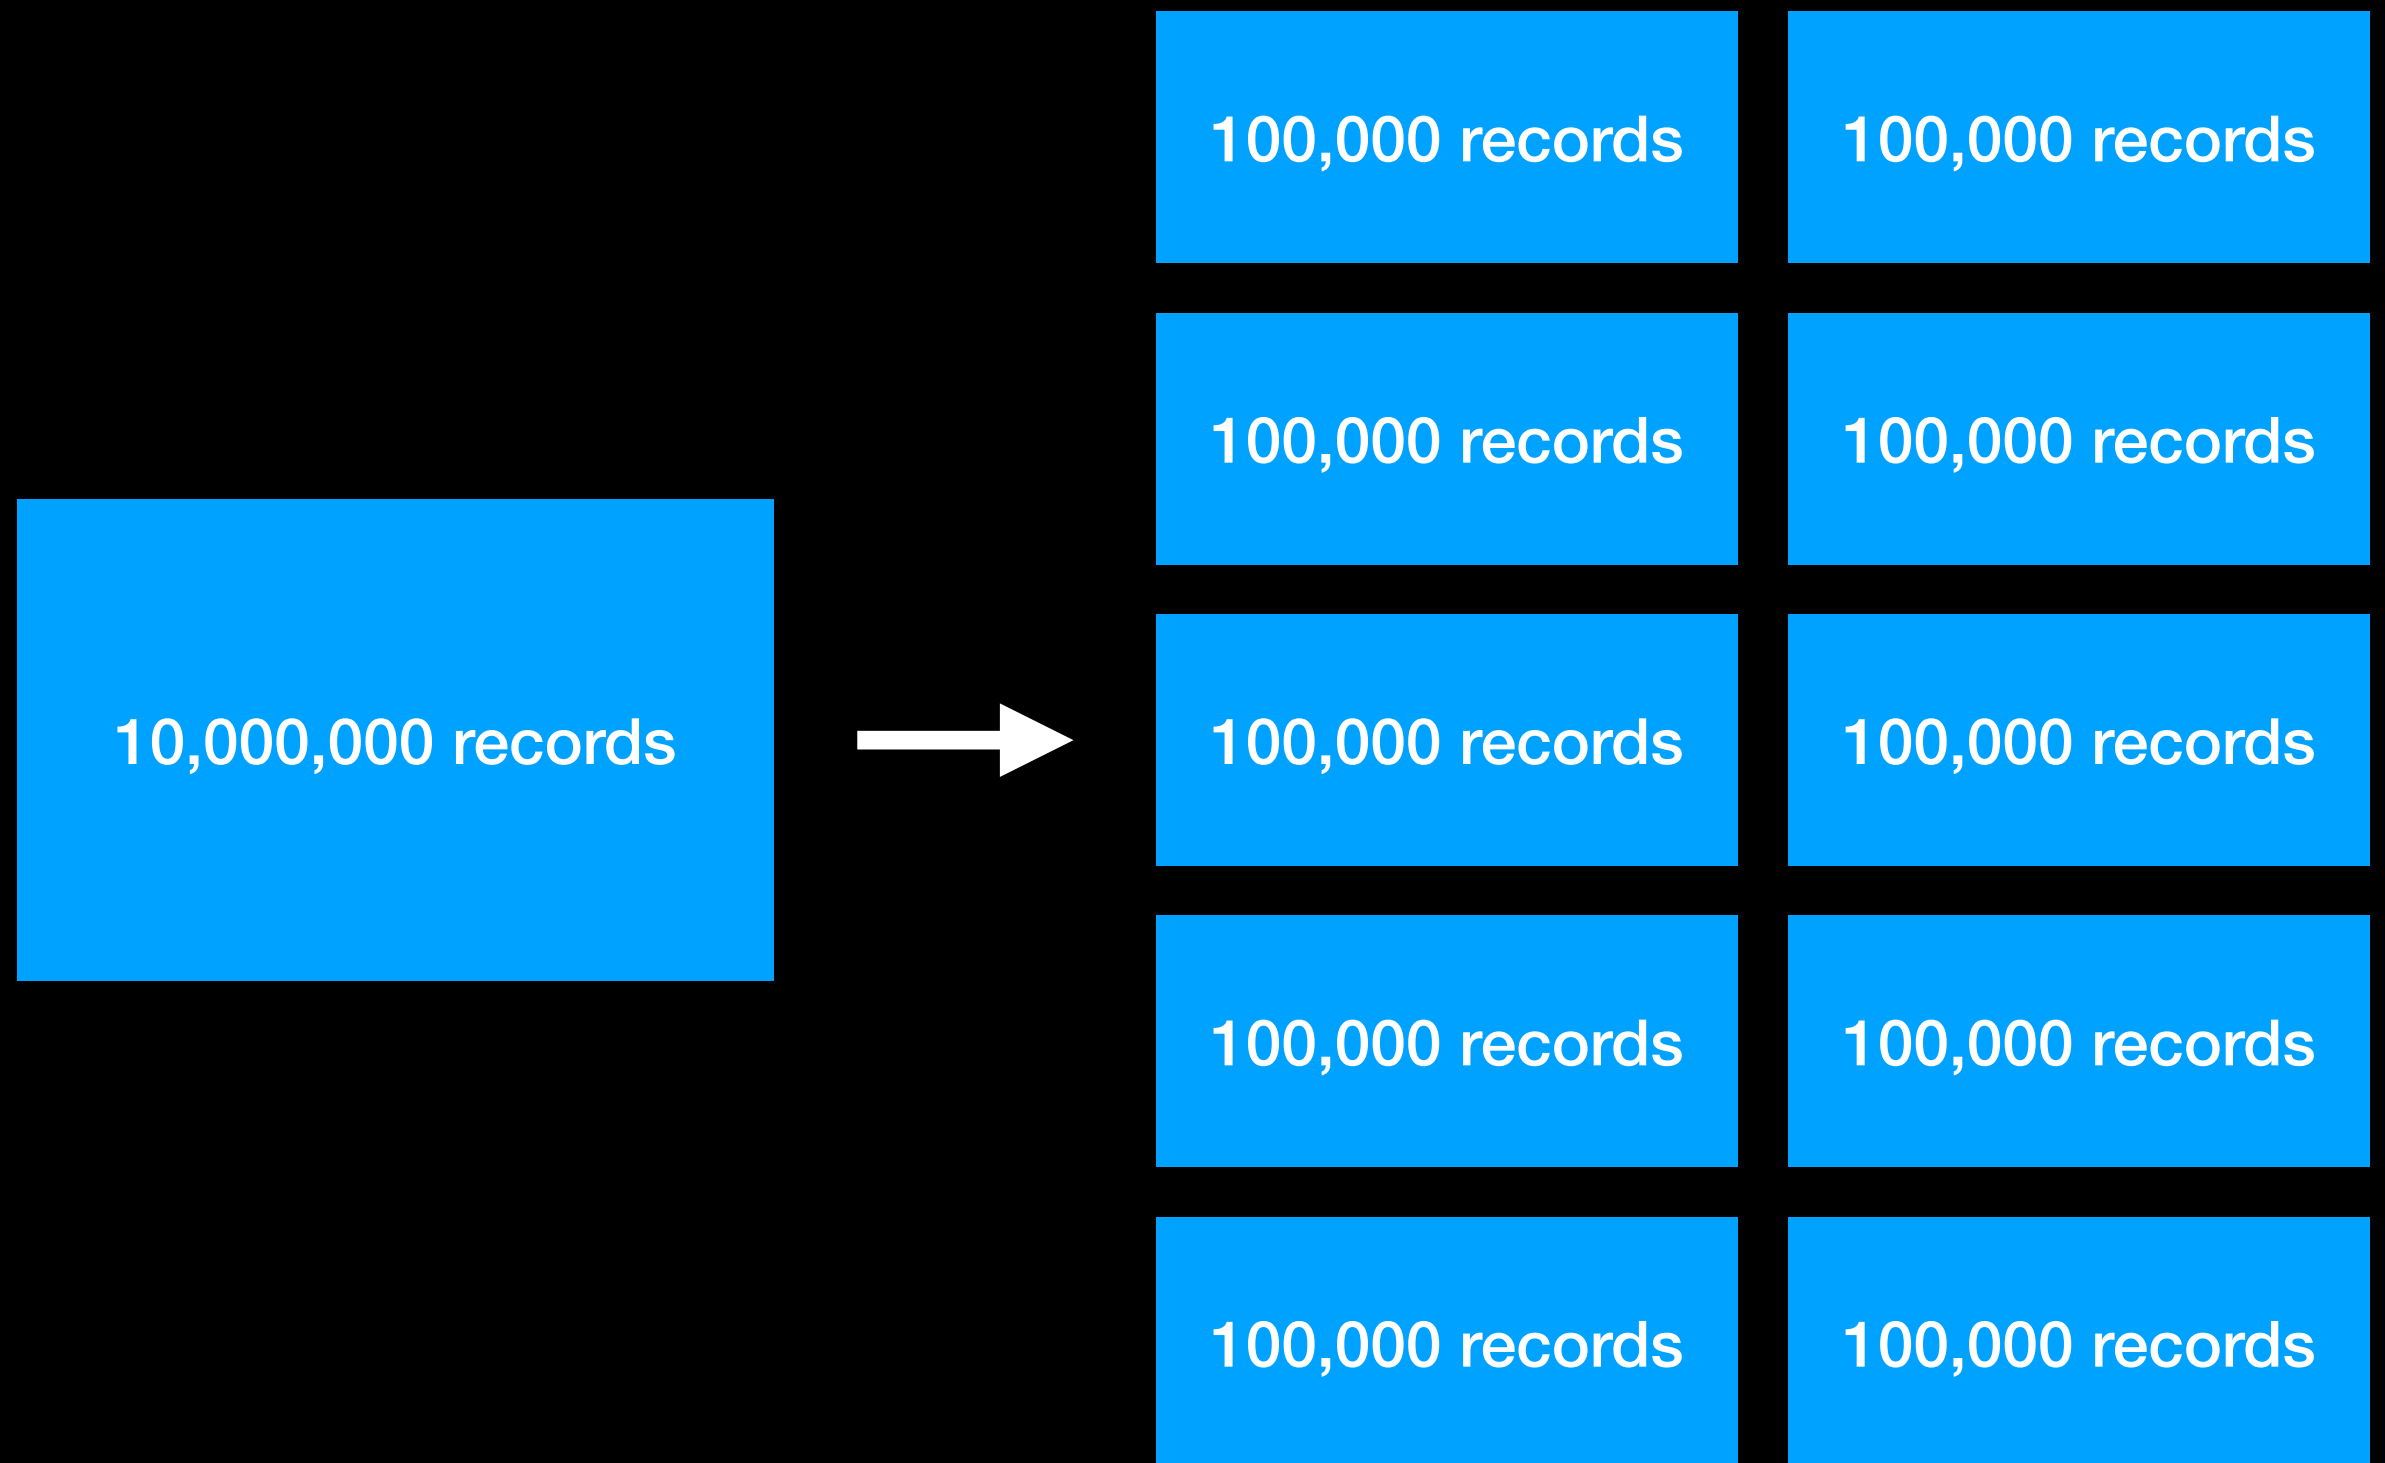
\$\$\$\$\$

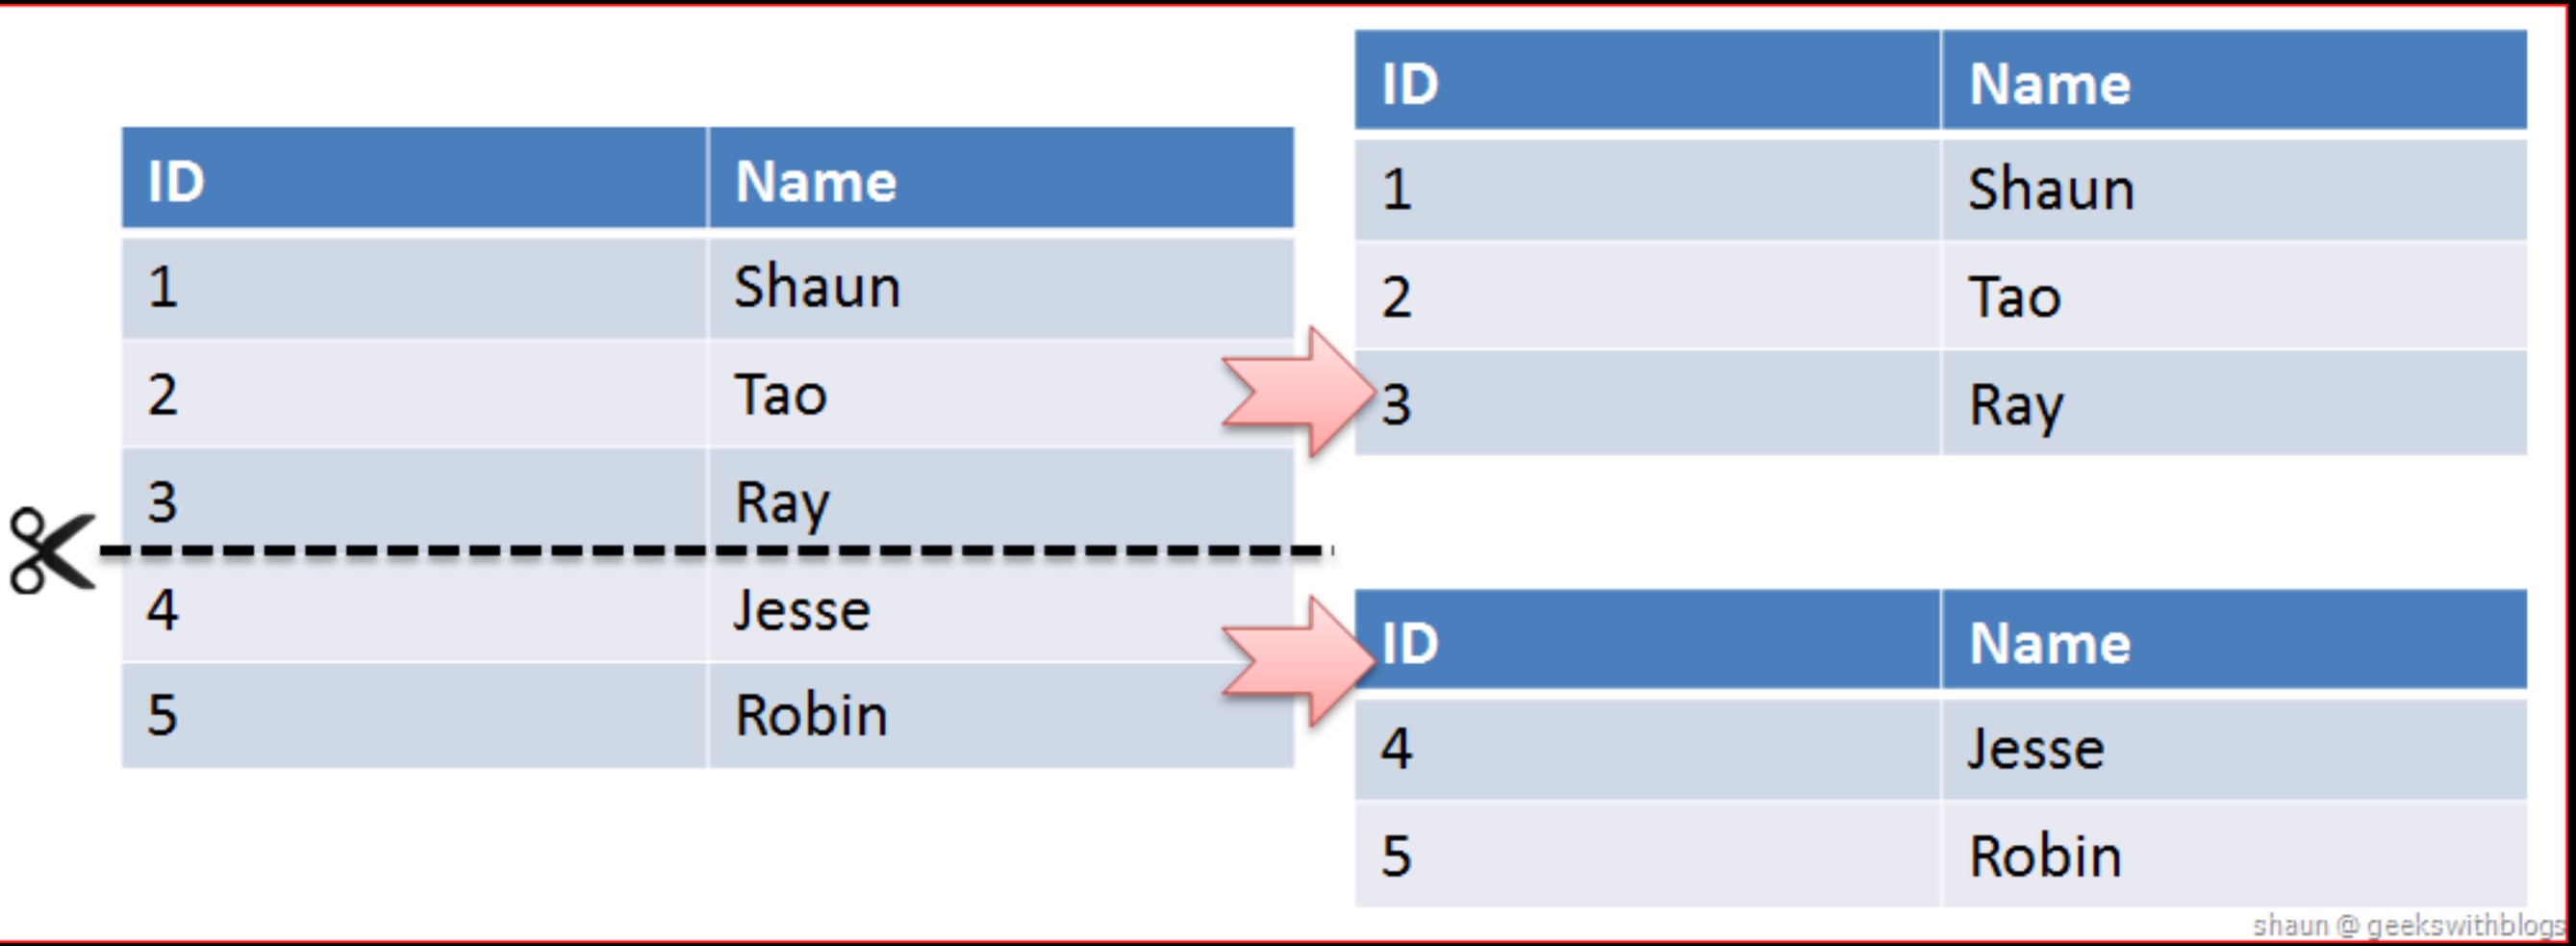
What's **Sharding**?

Vertical Partitioning

Horizontal Partitioning

10,000,000 records





Sounds Easy?

Then, How to **Shard** ?

Shards

app_db

Server #1

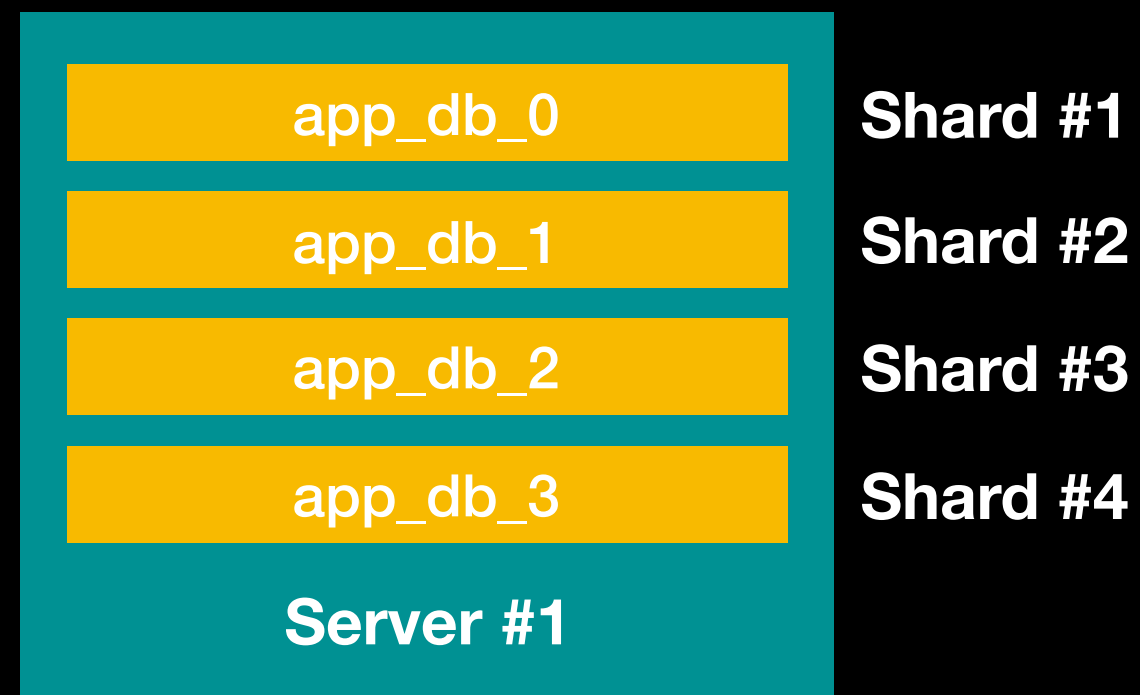
app_db_0

app_db_1

app_db_2

app_db_3

Server #1



app_db_0

app_db_1

app_db_2

app_db_3


Server #1

Shard #1 (user id 0~100)

Shard #2 (user id 101~200)

Shard #3 (user id 201~300)

Shard #4 (user id 301~400)



Multiple Shards per Server

app_db_0

Server #1

Shard #1

app_db_1

Server #2

Shard #2

app_db_2

Server #3

Shard #3

app_db_3

Server #4

Shard #4

One Shard per Server

Shard Key

Shard Key is the key
you use to split the data

Key

Key -> Compute Shard

-> Shard #

user_id -> hash -> shard #5

store_id -> hash -> shard #2

Types of Shard Keys

- Hash
- Range
- User-defined (Tags in MongoDB)

How to choose the
Shard Key ?

It depends on how you
find / query the data

Example #1

- Company (id)
- Store (id, company_id)
- Order (id, store_id, amount ...)
- OrderItem (order_id, quantity, ...)
- OrderPayment (order_id...)

Then you need to analysis
across different stores

```
SELECT SUM(amount) FROM orders  
WHERE store_id IN (2,3,4,5)  
GROUP BY store_id
```


Your key is
the ID of the stores

Example #1

- Company (id)
- Store (id, company_id)
- Order (id, **company_id**, **store_id**, amount ...)
- OrderItem (**company_id**, **store_id**, order_id, quantity, ...)
- OrderPayment (**company_id**, **store_id**, order_id ...)

And then you can split
the tables by **store_id**

The **store_id** could also
be a part of the **UUID** of
the order

And so you can extract the
store_id from the **UUID** to
lookup the **shard**

Different Sharding Approaches

- Database level sharding
- SQL level sharding
- Middleware-based sharding
- Application level sharding

RDBMS Solutions

- Vitas (from YouTube, Written in Go)
- MySQL Cluster (NDB)
- MySQL Fabric (Middleware-based sharding in Python)
- MariaDB Spider Engine
- GPDB (PostgreSQL)

Vitas

- Written in Go
- Use sql parser to distribute the queries
- Only for MySQL
- protocol gRPC
- Used by YouTube

MySQL Cluster

- Limited to 48 hosts
- Shard Key MUST BE in primary key and can't be changed. Similar to MySQL partitions.
- Primary key could be compound primary key
- NDB Engine
- When query doesn't include shard key, query will be scattered query.

MySQL Fabric

- Middleware-based sharding
- Written in Python
- Has it's own connector API binding for different languages (Java, Python)
- Manually set the host to operate from the application code.
- Bottleneck will be the middleware.

MySQL 上 sharding 的方案

Posted on [February 2, 2017](#) by [Gea-Suan Lin](#)

 Like 14

  0

[Percona](#) 的人剛好針對 database sharding 的事情整理了一篇文章，專門講 SaaS 服務時對 sharding 的規劃：「[MySQL Sharding Models for SaaS Applications](#)」。

我主要是看這段：


When sharding today, you have a choice of rolling your own system from scratch, using comprehensive sharding platform such as [Vitess](#) or using a proxy solution to assist you with sharding. For proxy solutions, [MySQL Router](#) is the official solution. But in reality, third party solutions such as open source [ProxySQL](#), commercial [ScaleArc](#) and semi-commercial (BSL) [MariaDB MaxScale](#) are widely used. Keep in mind, however, that traffic routing is only one of the problems that exist in large scale sharding implementations.


也就是說，除了在 application 上自己實作外，目前能用的方案(堪用的方案)大概就這些了，而且這些方案主要是解 routing 問題，對於跨 database server 的 join 都還是要自己小心實作。


by @gslin

Consistent Hashing

2^{32} | 0










● User(UUID=7fd7ed08-5da8-11e7-9bb5-3c15c2cb5a5a)

● User(UUID=7fd7ed08-5da8-11e7-9bb5-3c15c2cb5a5a)







PHP **CRC32** returns 2^{32}
positive integer on 64bit
machine

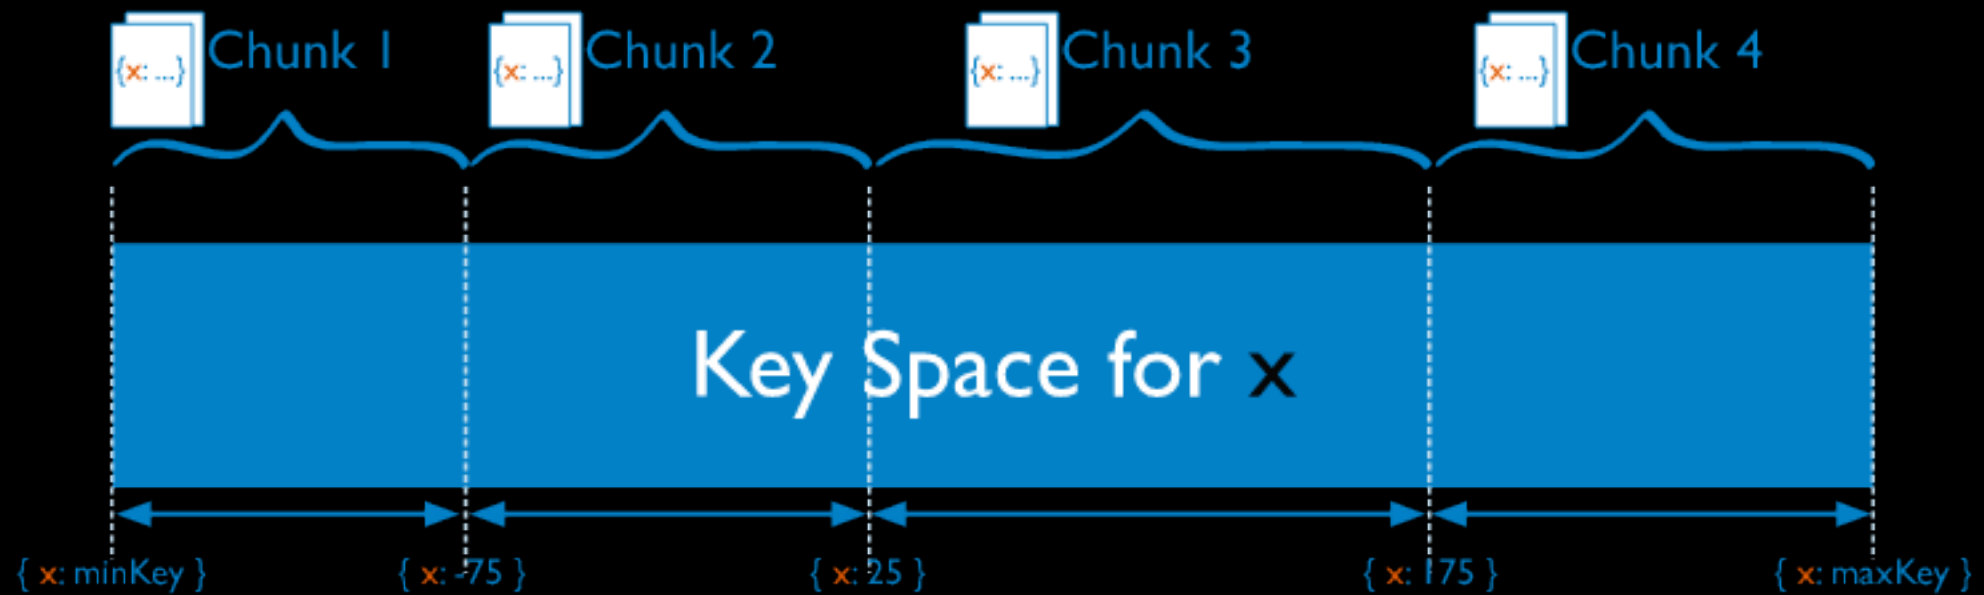
And so we use the **native
PHP array** to record them

```
foreach ($mapping->chunks as $i => $c) {  
    $x = $c['index'];  
    $this->buckets[$x] = $i;  
    $this->targetIndexes[$i][] = $x;  
}  
ksort($this->buckets, SORT_REGULAR);
```

Chunks

Split the key space of the
shard key into **chunks**

Chunks



And each **chunk** is
assigned to a **shard**

- Chunk #1
- Chunk #2
- Chunk #3
- Chunk #4
-

- Key(0~100) Chunk #1
- Key(101~200) Chunk #2
- Key(201~300) Chunk #3
- Key(301~400) Chunk #4
-

- Key(0~100) Chunk #1 -> Shard #1
- Key(101~200) Chunk #2 -> Shard #1
- Key(201~300) Chunk #3 -> Shard #2
- Key(301~400) Chunk #4 -> Shard #2
-

```
$mapping = new ShardMapping('store_keyspace', [
    'key' => 'store_id',
    'shards' => ['node1', 'node2', 'node3'],
    'chunks' => [
        ["from" => 0, "index" => 536870912, "shard" => "node1" ],
        ["from" => 536870912, "index" => 1073741824, "shard" => "node1" ],
        ["from" => 1073741824, "index" => 1610612736, "shard" => "node1" ],
        ["from" => 1610612736, "index" => 2147483648, "shard" => "node2" ],
        ["from" => 2147483648, "index" => 2684354560, "shard" => "node2" ],
        ["from" => 2684354560, "index" => 3221225472, "shard" => "node2" ],
        ["from" => 3221225472, "index" => 3758096384, "shard" => "node3" ],
        ["from" => 3758096384, "index" => 4294967296, "shard" => "node3" ],
    ]
], $dataSourceManager);

$hasher = new FastHasher($mapping);
$hash = $hasher->hash($key); // return 3221225472
```

Shard Balancing

Shard Migration

- Server #1 (10 shards)
- Server #2 (10 shards)
- Server #3 (10 shards)

- Server #1 (10 shards)
- Server #2 (10 shards)
- Server #3 (10 shards)

Server load increases because server #1 has more rows than other servers.

- Server #1 (9 shards)
- Server #2 (10 shards)
- Server #3 (10 shards)
- Server #4 (1 shard)

Migrate the bottleneck shard from Server #1 to a new server.

The underlying tool
mysqldbcopy

```
mysqldbcopy --replication=slave  
--locking=lock-all  
--source=root:pass@host1  
--destination=root:pass@host2  
database
```

replication mode=slave: copy a database from one slave to another attached to the same master

```
[root@localhost /]# mysqldbcopy --source=superadmin:root@192.168.1.10:3306 --destination=superadmin:root@192.168.1.20:3306
--all --force --locking=snapshot --rpl master --rpl-user=repl:root
# Source on 192.168.1.10: ... connected.
# Destination on 192.168.1.20: ... connected.
# Including all databases.
# GTID operation: SET @MYSQLUTILS_TEMP_LOG_BIN = @@SESSION.SQL_LOG_BIN;
# GTID operation: SET @@SESSION.SQL_LOG_BIN = 0;
# GTID operation: SET @@GLOBAL.GTID_PURGED = 'f2e752b0-0a94-11e3-87dd-080027c9399e:1-608';
# Copying database db1
# Copying TABLE db1.foo
# Copying TABLE db1.gaga
# Copying TABLE db1.gtid
# Copying PROCEDURE db1.load_foo_test_data
# Copying PROCEDURE db1.one_load_foo_test_data
# Copying GRANTS from db1
# Copying database percona
# Copying TABLE percona.checksums
# Copying database test
# Copying TABLE test.heartbeat
# Copying GRANTS from test
# Copying data for TABLE db1.foo
# Copying data for TABLE db1.gaga
# Copying data for TABLE db1.gtid
# Copying data for TABLE percona.checksums
# Copying data for TABLE test.heartbeat
# GTID operation: SET @@SESSION.SQL_LOG_BIN = @MYSQLUTILS_TEMP_LOG_BIN;
# Connecting to the current server as master
#...done.
```

```
use Maghead\Sharding\Operations\CloneShard;
```

```
$op = new CloneShard($config);
```

```
// public function clone($mappingId, $instanceId, $newNodeId, $srcNodeId)
```

```
$op->clone("store_keyspace", "server3", "shard1000", "shard0003");
```

```
maghead shard clone
  --drop-first
  --mapping store_key
  --instance server2
a01 a11
```

clone shard "a01" to the a new shard on server2 named "a11"

```
maghead shard allocate  
  --mapping store_key  
  --instance server2  
  a02
```

allocate an empty shard "a02" on server 2 and initialize all the tables

Shard Balancing API

```
$balancer = new ShardBalancer(  
    new ConservativeShardBalancerPolicy(1, false, 1.3)  
);  
  
$migrateInfo = $balancer->balance($mapping,  
                                   [new OrderSchema]);
```

```
$balancer = new ShardBalancer(  
    new ConservativeShardBalancerPolicy(1, false, 1.3)  
);  
  
$migrateInfo = $balancer->balance($mapping,  
    [new OrderSchema]);
```

保守式平衡規則

```
$balancer = new ShardBalancer(  
    new ConservativeShardBalancerPolicy(1, false, 1.3)  
);  
  
$migrateInfo = $balancer->balance($mapping,  
    [new OrderSchema]);
```

一次只找出一個 Chunk 做 Migration

```
$balancer = new ShardBalancer(  
    new ConservativeShardBalancerPolicy(1, false, 1.3)  
);  
  
$migrateInfo = $balancer->balance($mapping,  
    [new OrderScheme]);
```

找出 rows 總量超過平均 1.3 倍的 Shard

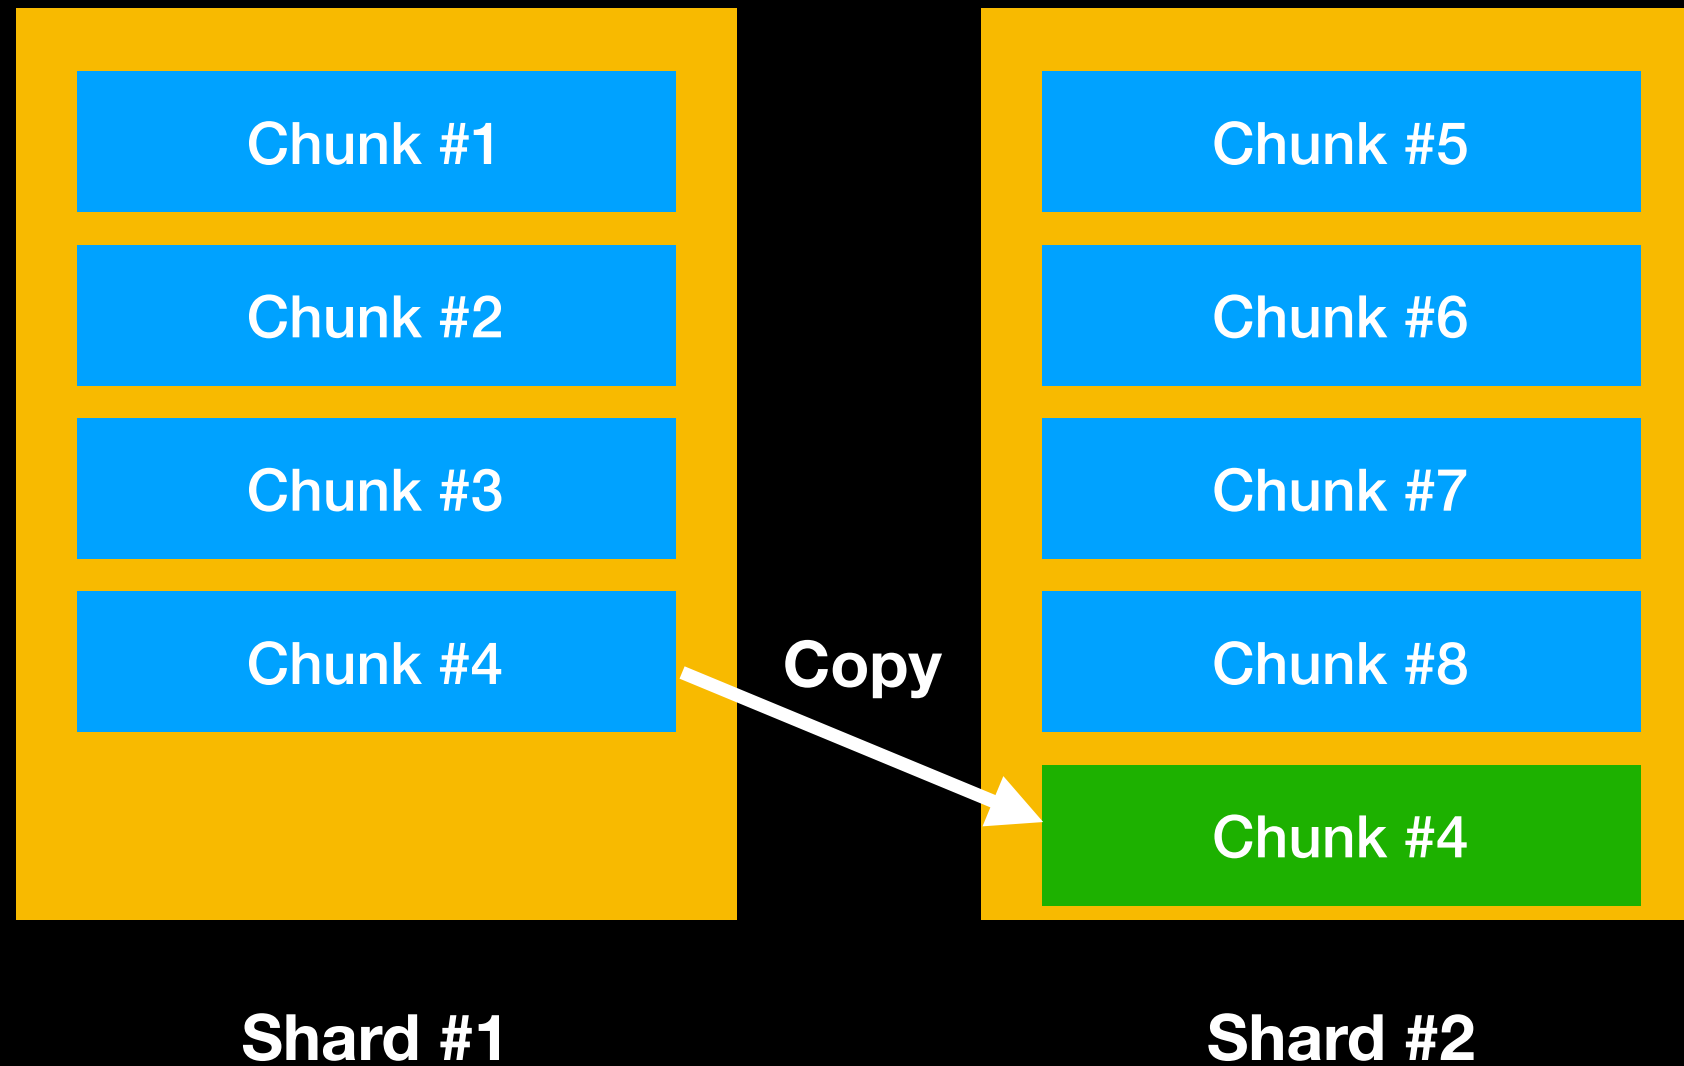
Chunk Migration

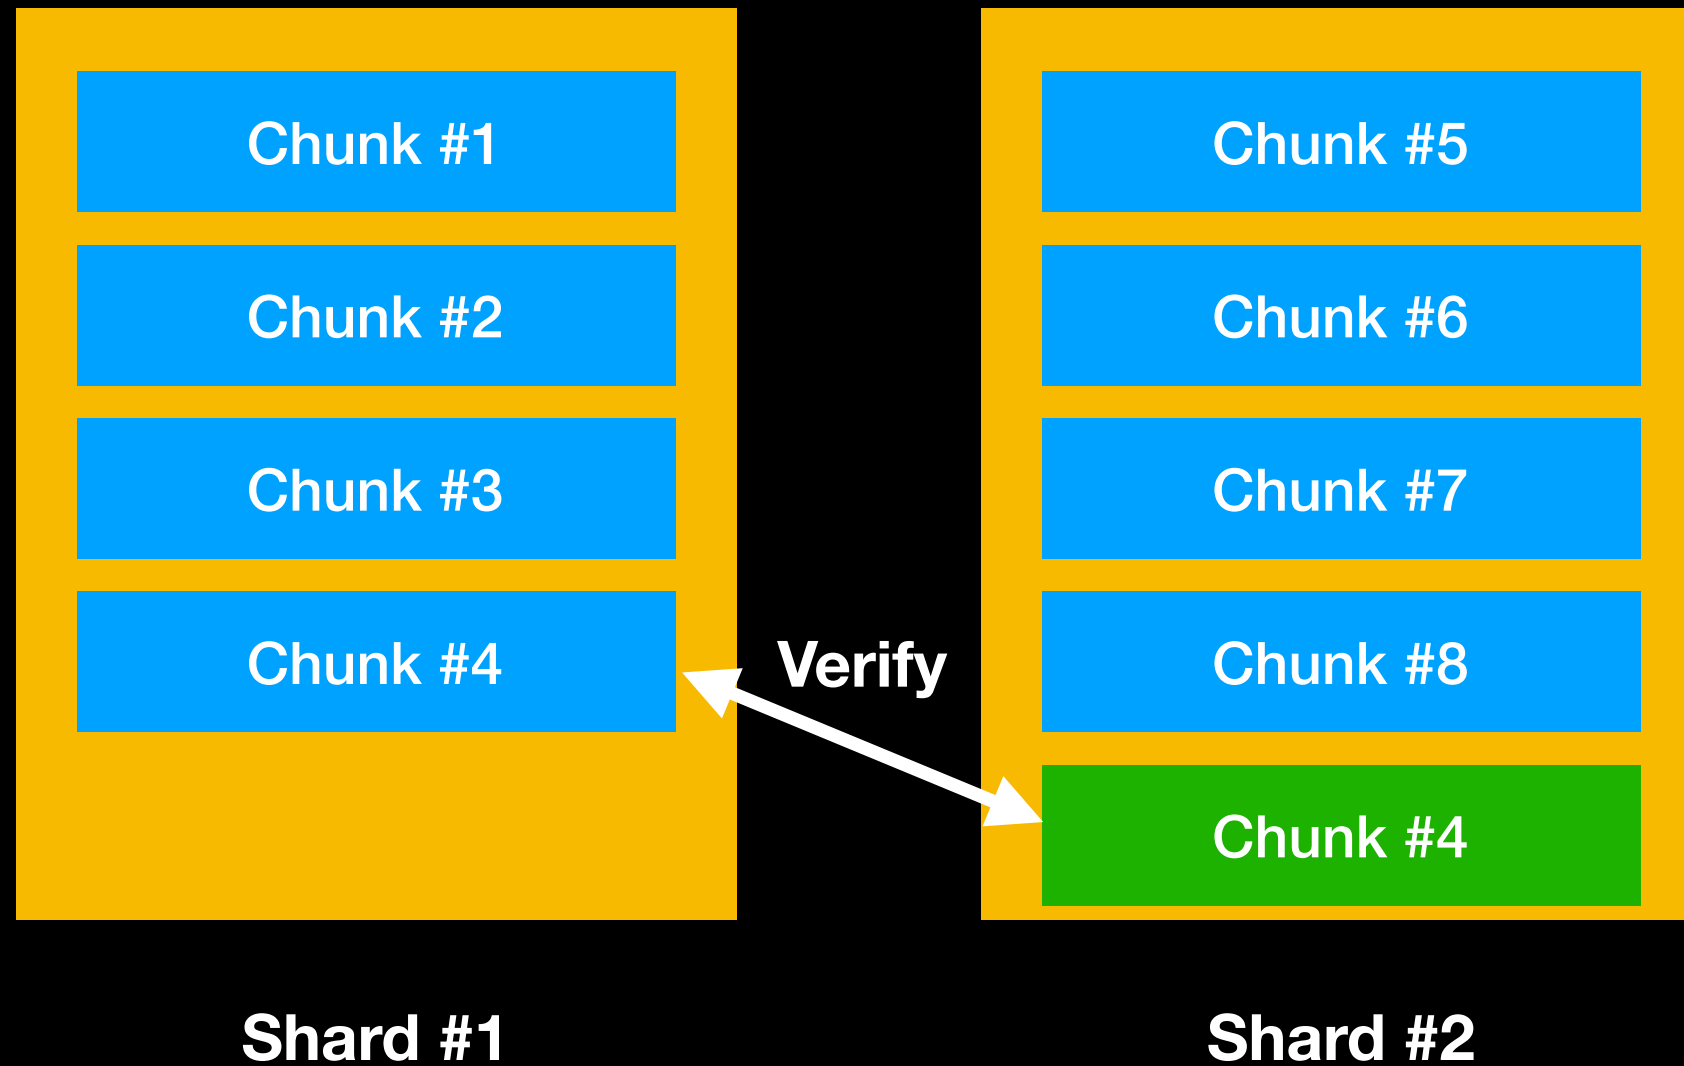
- Key(0~100) Chunk #1 -> Shard #1
- Key(101~200) Chunk #2 -> Shard #1
- Key(201~300) Chunk #3 -> Shard #2
- Key(301~400) Chunk #4 -> Shard #2
-

- Key(0~100) Chunk #1 -> Shard #1
- Key(101~200) Chunk #2 -> Shard #1
- Key(201~300) Chunk #3 -> Shard #2
- Key(
-

When the load of shard #2 increases, we can assign some chunks to the other shard.

- Key(0~100) Chunk #1 -> Shard #1
- Key(101~200) Chunk #2 -> Shard #1
- Key(201~300) Chunk #3 -> Shard #2
- Key(301~400) Chunk #4 -> Shard #2
-
 - When the rows in Chunk #4 grow up too much,
We can split the Chunk #4 into 2+ chunks and
assign them to the other shards.







Shard #1



Shard #2

ChunkManager API

```
namespace Maghead\Sharding\Manager;
```

```
class ChunkManager
```

```
{
```

```
    public function __construct(ShardMapping $mapping) { }
```

```
    public function distribute(array $shardIds, $numberOfChunks = 32) { }
```

```
    public function verify(Chunk $chunk, Shard $dstShard, array $schemas) { }
```

```
    public function remove(Chunk $chunk, array $schemas) { }
```

```
    public function clone(Chunk $chunk, Shard $dstShard, array $schemas) { }
```

```
    public function migrate(Chunk $chunk, Shard $dstShard, array $schemas) { }
```

```
    public function move(Chunk $chunk, Shard $dstShard, array $schemas) { }
```

```
    public function split(Chunk $chunk, $n = 2) { }
```

```
}
```

```
$result = $chunkManager->migrate($chunk2, $shards['node3'], $schemas);  
if (!$result->isSuccessful()) {  
    // ...  
}
```

Jumbo Chunk

Chunks contain a large
number of rows

Migrating Jumbo Chunks
requires more time

Migration should start
from the **small chunks**

And in the **same zone**

Start Sharding in **Maghead**

Define Shard Mapping

or Keyspace

maghead shard mapping add

--hash

-s node1

-s node2

-s node3

--key store_id

M_store_id

sharding:

mappings:

M_store_id:

key: "store_id"

shards:

- node1
- node2
- node3

hash: true

chunks:

- { from: 0, index: 536870912, shard: node1 }
- { from: 536870912, index: 1073741824, shard: node1 }
- { from: 1073741824, index: 1610612736, shard: node1 }
- { from: 1610612736, index: 2147483648, shard: node2 }
- { from: 2147483648, index: 2684354560, shard: node2 }
- { from: 2684354560, index: 3221225472, shard: node2 }
- { from: 3221225472, index: 3758096384, shard: node3 }
- { from: 3758096384, index: 4294967296, shard: node3 }

自動分派的 Chunk 定義


```
class StoreSchema extends DeclareSchema
{
    public function schema()
    {
        $this->globalTable("M_store_id");
    }
}
```

Store table 需要在 Shards 中被 Join 所以被定義為 Global Table

```
namespace StoreApp\Model;  
  
use Maghead\Schema\DeclareSchema;  
  
class OrderSchema extends DeclareSchema  
{  
    public function schema()  
    {  
        $this->column('uuid', 'uuid-pk');  
  
        {...cut....}  
  
        $this->shardBy("M_store_id");  
    }  
}
```

Order table 需要依照定義的 keyspace 做分片

```
namespace StoreApp\Model;

use Maghead\Schema\DeclareSchema;

class OrderSchema extends DeclareSchema
{
    public function schema()
    {
        $this->column('uuid', 'uuid-pk');

        {....cut.
        $this->shardBy("M_store_id");
    }
}
```

快速定義 UUID primary key

maghead schema build

Use the model as usual

```
$ret = Order::create([  
    'store_id' => $storeId,  
    'amount' => 600,  
]);
```

Automatically dispatch record to the shard by the store_id

```
$ret->shard; // Shard object
```



The dispatched shard is returned in the result object.

Selecting Shards


```
$store = Store::masterRepo()->findByCode('TW002');
```

```
$shard = Order::shards()->dispatch($store->id);
```

```
// Maghead\Sharding\Shard
```

Deletion works as well

```
$order = Order::findByPrimaryKey($key);
```

```
$ret = $order->delete();
```

ShardCollection

```
$shards = Order::shards();  
$shards->locateBy(function($repo, $shard) {  
    return $repo->findByCode('X01');  
});
```

```
$shards = Order::shards(); // ShardCollection  
$shards->map(function($repo, $shard) {  
    // do something  
});
```

QueryWorker

- Maghead\Sharding\QueryMapper\GearmanQueryMapper
- Maghead\Sharding\QueryMapper\PthreadQueryMapper


```
$query = new SelectQuery;  
$query->select(['SUM(amount)' => 'amount']);  
$query->from('orders');  
$query->where()->in('store_id', [1,2]);
```

```
$client = new GearmanClient;  
$client->addServer();  
$mapper = new GearmanQueryMapper($client);  
$res = $mapper->map($shards, $query);
```

Reducer PHP7 Extension

Data Group By implemented in PHP Runtime

```
$result = group_by($rows, $fields, [  
    'total_amount' => [  
        'selector' => 'amount',  
        'aggregator' => REDUCER_AGGR_SUM,  
    ],  
    'cnt' => REDUCER_AGGR_COUNT,  
]);  
print_r($result);
```

```
$rows = [  
    [ 'category' => 'Food', 'type' => 'pasta', 'amount' => 1 ],  
    [ 'category' => 'Food', 'type' => 'pasta', 'amount' => 1 ],  
    [ 'category' => 'Food', 'type' => 'juice', 'amount' => 1 ],  
    [ 'category' => 'Food', 'type' => 'juice', 'amount' => 1 ],  
    [ 'category' => 'Book', 'type' => 'programming', 'amount' => 5 ],  
    [ 'category' => 'Book', 'type' => 'programming', 'amount' => 2 ],  
    [ 'category' => 'Book', 'type' => 'cooking', 'amount' => 6 ],  
    [ 'category' => 'Book', 'type' => 'cooking', 'amount' => 2 ],  
];  
$result = group_by($rows, ['category', 'type'], [  
    'total_amount' => [  
        'selector' => 'amount',  
        'aggregator' => REDUCER_AGGR_SUM,  
    ],  
    'cnt' => REDUCER_AGGR_COUNT,  
]);  
print_r($result);
```

See more on GitHub

Thank You

Join Us!

<https://github.com/maghead/maghead>

Follow me on twitter: @c9s